Safe module unloading in Experimental Oberon

Andreas Pirklbauer

1.10.2018

In the Oberon system[1], there exist three possible types of references to a loaded module M[2]:

1. *Client references* exist when other loaded modules *import* module M.

2. *Type references* exist when type tags (addresses of type descriptors) in *dynamic* objects reachable by other loaded modules refer to descriptors of types *declared* in module M.

3. *Procedure variable references* exist when procedure variables in *static* or *dynamic* objects reachable by other loaded modules refer to procedures *declared* in module M.

In most Oberon implementations, only *client* references are checked prior to module unloading. *Type* and *procedure variable* references are usually not checked, although various approaches are typically employed to address the case where modules *with* such references are unloaded[3].

In Experimental Oberon[4], *all* possible types of references to a loaded module or module group are checked prior to module *unloading*, which is implemented as follows (see Figure 1):

- If clients exist, a module or module group is never unloaded.
- If no client, type or procedure variable references to a module or module group exist in the remaining modules or data structures, it is unloaded and its associated memory is released.
- If no clients, but type or procedure variable references exist, there are two possibilities: (a) If no option is specified in the module *unload* command, the command takes no action and merely displays the names of the modules containing the references that caused the removal to fail. (b) If the *force* option is specified, the specified modules are initially removed only from the *list* of loaded modules, without releasing their associated memory[5]. Such "hidden" modules are later removed from *memory* using the command *Modules.Collect*, as soon as there are no more references to them from anywhere in the system.

---

[1] *http://www.projectoberon.com*
[2] *An Oberon module can be viewed as a container of types, variables and procedures, where variables can be procedure-typed. Types can be declared as g l o b a l types (in which case they can be exported and referenced by name in clients) or as types l o c a l to a procedure (in which case they cannot be exported). Variables can be statically declared as g l o b a l variables (allocated in the module area when a module is loaded) or as l o c a l variables (allocated on the stack when a procedure is called), or they can be dynamically allocated in the h e a p via the predefined procedure NEW. Procedures can be declared as global or local procedures. They can also be assigned to g l o b a l procedure variables (of the same or a different module) or to procedure variable fields in d y n a m i c objects – even if they are not exported. Thus, in general there can be type, variable, procedure and procedure variable references from both static and dynamic objects of other modules to static or dynamic objects of the modules to be unloaded. However, only d y n a m i c t y p e references and s t a t i c and d y n a m i c p r o c e d u r e v a r i a b l e references need to be checked during module unloading for the following reasons: First, s t a t i c type, variable and procedure references from other loaded modules can only refer by n a m e to types, variables or procedures d e c l a r e d in the modules to be unloaded. Such references are already handled via their import/export relationship during module unloading (if clients exist, a module or module group is never unloaded) and therefore don't need to be checked separately. Second, d y n a m i c pointer variable references from global or dynamic p o i n t e r variables of other modules to d y n a m i c objects reachable by the modules to be unloaded s h o u l d not be checked, as such references s h o u l d not prevent module unloading. In the Oberon system, such references will be handled by the garbage collector during a future garbage collection cycle, i.e. heap records reachable by the just unloaded modules a n d other still loaded modules will not be collected, whereas heap records that w e r e reachable o n l y by the just unloaded modules w i l l be collected – as they should. Thus, the handling of dynamic pointer references is delegated to the garbage collector. Finally, p o i n t e r variable references to s t a t i c objects declared as global variables are only possible by resorting to low-level facilities and should be avoided (and, in fact, be disallowed).*
[3] *See the appendix for historical notes on module unloading in the Oberon system.*
[4] *http://www.github.com/andreaspirklbauer/Oberon-experimental*
[5] *The force option /f must be specified at the e n d of the list of modules to be unloaded, e.g., System.Free M1 M2 M3/f*
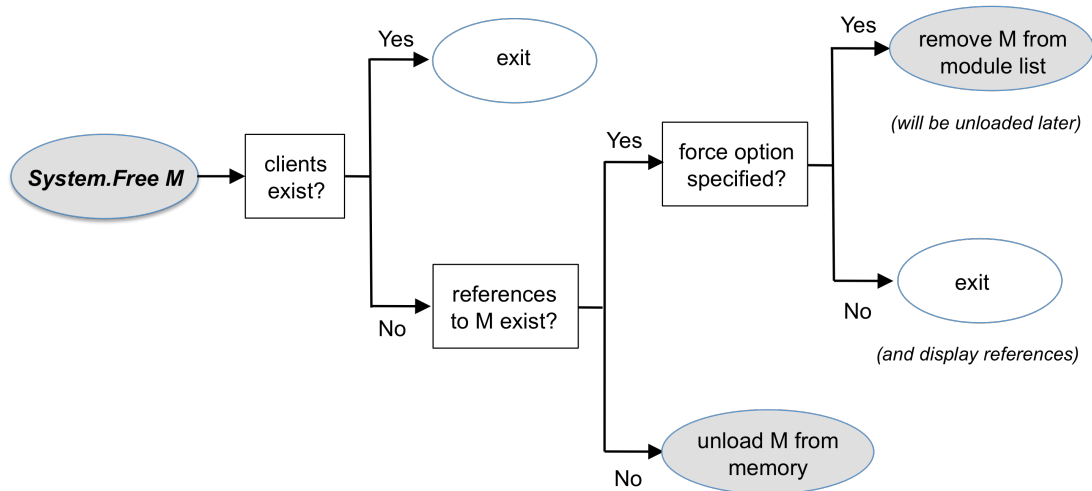
Figure 1: Safe module unloading in Experimental Oberon

Removing modules only from the *list* of loaded modules amounts to *renaming* them, allowing other modules with the same names to be reloaded again, without having to unload (from memory) earlier versions that are still referenced[6]. Removing a module from *memory* frees up the memory area previously occupied by the module block[7].

To make the removal of no longer referenced hidden module data *automatic*, the command *Modules.Collect* is included in the Oberon background task handling garbage collection. It checks all possible combinations of k modules chosen from n hidden modules for clients and references, and removes those module subgroups from memory that are no longer referenced. The command *Modules.Collect* can also be manually activated at any time. Alternatively, the user can invoke *System.Collect*, which includes a call to *Modules.Collect*.

In sum, module unloading does not affect past references, as module data is kept in memory for exactly as long as necessary and removed from it as soon as possible.

For example, older versions of a module's code can still be executed if they are referenced by static or dynamic procedure variables in other modules, even if a newer version of the module has been reloaded in the meantime[8]. Type descriptors also remain accessible to other modules for exactly as long as needed. This covers the important case where a structure rooted in a variable of base type T declared in a base module M contains elements of an extension T' defined in a client module M', which is unloaded[9]. Such elements typically contain both *type* references (type tags) and *procedure variable* references (handler procedures) referring to M'. This is common in the Oberon viewer system, for example, where M is module *Viewers*.

---

[6] *Modules removed only from the list of loaded modules, but not from memory, are marked with an asterisk in the output of the command System.ShowModules. Commands of such "hidden" modules can be accessed by either specifying their module number or their (modified) module name, both of which are displayed by the command System.ShowModules. In both cases, the corresponding command text must be enclosed in double quotes. If a module M carries module number 14, for instance, one can activate a command M.P also by clicking on the text "14.P". Typical use cases include hidden modules that still have background tasks installed which can only be removed by activating a command of the hidden modules themselves, or hidden modules that still have open viewers. If the command to close a viewer is displayed in the viewer's menu bar, the user can manually edit the command text (by clicking within its bottom two pixel lines), replace the module name by its module number and enclose the modified command text in double quotes. Although this is somewhat clumsy, it at least enables the user to close the viewer. An alternative approach is to provide a "Close" command that also accepts the marked viewer as argument (using procedure Oberon.MarkedViewer).*

[7] *In Original Oberon 2013 on RISC and in Experimental Oberon, the module block includes the module's type descriptors. In some other Oberon implementations, such as Oberon on Ceres, type descriptors are not stored in the module block, but are dynamically allocated in the heap at module load time, in order to persist them beyond the lifetime of their associated modules. In Experimental Oberon, no such special precaution is necessary, as module blocks are removed only from the list of loaded modules and not from memory, if they are still referenced by other modules. Thus, type descriptors can safely be stored in the (static) module blocks.*

[8] *If an older version of a module's code accesses global variables (of itself or of other modules), it will automatically access the "right" version of such variables – as it should.*

[9] *Note that for t y p e references, it is actually possible to know at compile time, whether a module M may potentially lead to references from other modules at run time (but not whether it actually will): If M does not declare record types which are extensions T' of an imported type T, records declared in M cannot be inserted in a data structure rooted in a variable of an imported type T – precisely because they are not extensions of T (in the Oberon language, the assignment p := p' is allowed only if the type of p' is the same as the type of p or an extension of it).*

If a module *group* is to be unloaded and there exist clients or references *only* within this group, it is unloaded *as a whole*. This can be used to remove module groups with *cyclic* references[10].

It is also possible to release *entire subsystems* of modules. The command *System.FreeImports* attempts to unload the specified modules and all their direct and indirect *imports*. It may be used for conventional module stacks with a single top module, for example a compiler. By contrast, the command *System.FreeClients* unloads the specified modules and all their direct and indirect *client*s. This variant may be used for module stacks written in the object-oriented programming style with their typically inverted module hierarchy, for example a graphics editor.

Note that these unloading strategies amount to *heuristics* and may tend to unload rather large parts of the system, which may not be desired. Thus, the recommended way to unload modules is to use the command *System.Free* with a *specific* set of modules provided as parameters.

For added convenience, the tool commands *System.ShowRefs* and *System.ShowExternalRefs* can be used to identify all modules containing references to a given module or module group.

Implementation aspects

When the user attempts to unload a module or module group, the unload command first *selects* the modules to be unloaded using procedure *Modules.Select* and then invokes procedure *Modules.Check*. Clients are checked first. This is accomplished by simply verifying whether *unselected* modules import *selected* modules[11]:

```
1   PROCEDURE FindClients*(proc: ModHandler; VAR res: INTEGER);
2     VAR mod, imp, prev, m: Module; p, q: INTEGER; continue: BOOLEAN;
3   BEGIN res := 0; mod := root; continue := proc # NIL; prev := NIL;
4     WHILE continue & (mod # NIL) DO
5       IF (mod.name[0] # 0X) & mod.selected & (mod.refcnt > 0) THEN m := root;
6         WHILE continue & (m # NIL) DO
7           IF (m.name[0] # 0X) & ~m.selected THEN p := m.imp; q := m.cmd;
8             WHILE p < q DO imp := Mem[p];
9               IF imp = mod THEN INC(res, proc(m, mod, prev, continue)); prev := mod; p := q ELSE INC(p, 4) END
10            END
11          END ;
12          m := m.next
13        END
14      END ;
15      mod := mod.next
16    END
17  END FindClients;
```

If clients exist, no further action is taken and the module unload command exits.

If no clients exist, references are checked next. As mentioned earlier, references that *need* to be checked prior to module unloading include *type tags* (addresses of type descriptors) in *dynamic* heap objects reachable by all other loaded modules pointing to descriptors of types declared in the modules to be unloaded, and *procedure variables* installed in *dynamic or static* objects of other modules referring to *static* procedures declared in the modules to be unloaded.

---

[10] *In Oberon, cyclic references can be created by pointers or procedure variables, whereas cyclic module imports are normally not allowed. However, through a tricky sequence of compilation and editing steps, it is in fact possible to c o n s t r u c t cyclic module imports, which cannot be detected by the compiler. But even though they can be created, such modules are not allowed to be loaded, as the module loader of Experimental Oberon – adopting the approach chosen in Original Oberon – would simply enter an endless recursion, eventually leading to a stack overflow or out-of-memory condition – a totally acceptable solution for such an artificially constructed case. But even if modules with cyclic module imports w e r e allowed to be loaded, Experimental Oberon would handle them correctly upon unloading, i.e. if no external clients or references exist, such a module group would simply be unloaded as a whole – as it should.*
[11] *Mem stands for the entire memory and assignments involving Mem are expressed as SYSTEM.GET(a, x) for x := Mem[a] and SYSTEM.PUT(a, x) for Mem[a] := x.*

References from *dynamic* objects are checked using a conventional *mark-scan* scheme:

```
1    PROCEDURE FindDynamicRefs*(type, proc: RefHandler; VAR resType, resProc: INTEGER; all: BOOLEAN);
2     VAR mod: Module;
3    BEGIN mod := root;
4     WHILE mod # NIL DO
5      IF (mod.name[0] # 0X) & ~mod.selected THEN Kernel.Mark(mod.ptr);
6       IF ~all THEN Kernel.Scan(type, proc, mod.name, resType, resProc) END
7      END ;
8       mod := mod.next
9      END ;
10     IF all THEN Kernel.Scan(type, proc, "", resType, resProc) END
11    END FindDynamicRefs;
```

During the *mark* phase, heap records reachable by all *named* global pointer variables of *all other* loaded modules are marked (line 5), thereby excluding records reachable *only* by the specified modules themselves. This ensures that when a module *or module group* is referenced *only* by itself, it can still be unloaded. The *scan phase* (line 6 or 10), implemented as a separate procedure *Scan*[12] in module *Kernel*, scans the heap sequentially, unmarks all *marked* records and checks whether the *type tags* of the marked records point to descriptors of types in the *selected* modules to be unloaded, and also whether *procedure variables* declared in these heap records refer to procedures declared in those same modules.

An additional boolean parameter *all* allows the caller to specify whether the *mark* phase should first mark the heap records reachable by *all* other loaded modules before entering the *scan* phase (can be used for module unloading), or whether the *mark* phase should mark *each* of the other modules *individually* before initiating the *scan* phase (can be used for enumerating the references to each module *individually*).

Finally, the check for procedure variable references is also performed for all *global* procedure variables, whose offsets in the module's data section are obtained from an array in the module's *meta* data section, headed by the link *mod.pvr* in the module descriptor (see below):

```
1    PROCEDURE FindStaticRefs*(proc: RefHandler; VAR res: INTEGER);
2     VAR mod: Module; pref, pvadr, r: LONGINT; continue: BOOLEAN;
3    BEGIN res := 0; mod := root; continue := proc # NIL;
4     WHILE continue & (mod # NIL) DO
5      IF (mod.name[0] # 0X) & ~mod.selected THEN
6       pref := mod.pvr; pvadr := Mem[pref];
7       WHILE continue & (pvadr # 0) DO r := Mem[pref]
8        INC(res, proc(r, mod.name, continue));
9        INC(pref, 4); pvadr := Mem[pref]
10       END
11      END ;
12      mod := mod.next
13     END
14    END FindStaticRefs;
```

Note that the procedures *FindClients*, *FindDynamicRefs* and *FindStaticRefs* are expressed as *generic* traversal schemes which accept *parametric* handler procedures that are called for each encountered object. This allows these procedures to be used for *other* purposes as well, for example to *enumerate* or *count* the clients or references to a given module or module group.

---

[12] *The original procedure Kernel.Scan (implementing the scan phase of the Oberon garbage collector) has been renamed to Kernel.Collect, in analogy to procedure Modules.Collect.*

In order to omit in module *Kernel* any reference to the module list rooted in module *Modules*, procedure *Kernel.Scan* is also expressed as a *generic* heap scan scheme[13]:

```
1    PROCEDURE Scan*(type, proc: Handler; s: ARRAY OF CHAR; VAR resType, resProc: INTEGER);
2      VAR offadr, offset, p, r, mark, tag, size: LONGINT; continue: BOOLEAN;
3    BEGIN p := heapOrg; resType := 0; resProc := 0; continue := (type # NIL) OR (proc # NIL);
4      REPEAT mark := Mem[p+4];
5        IF mark < 0 THEN (*free*) size := Mem[p]
6        ELSE (*allocated*) tag := Mem[p]; size := Mem[tag];
7          IF mark > 0 THEN (*marked*) Mem[p+4] := 0;
8            IF continue THEN
9              IF type # NIL THEN INC(resType, type(tag, s, continue)) END ;
10             IF continue & (proc # NIL) THEN offadr := tag + 16; offset := Mem[offadr];
11               WHILE offset # -1 DO (*pointers*) INC(offadr, 4); offset := Mem[offadr] END ;
12               INC(offadr, 4); offset := Mem[offadr];
13               WHILE continue & (offset # -1) DO (*procedures*) r := Mem[p+8+offset];
14                 INC(resProc, proc(r, s, continue));
15                 INC(offadr, 4); offset := Mem[offadr]
16               END
17             END
18           END
19         END
20       END ;
21       INC(p, size)
22     UNTIL p >= heapLim
23   END Scan;
```

This scheme calls *parametric* handler procedures for individual elements of each *marked* heap record rather than *directly* checking whether these records contain *type* or *procedure variable* references to the modules to be unloaded. Procedure *type* is called with the *type tag* of the heap record as argument (line 9), while procedure *proc* is called for each procedure variable declared in the same record with (the address of) the *procedure* itself as argument (line 14). The results of the handler calls are *separately* added up for each handler and returned in the variable parameters *resType* and *resProc*.

An additional variable parameter *continue* allows the handler procedures to indicate to the caller that they are no longer to be called (lines 8, 10, 13). The scan process itself continues, but only to *unmark* the remaining marked records (line 7).

Procedure *Modules.Check* uses procedures *FindClients, FindDynamicRefs* and *FindStaticRefs* by passing its private handler procedures *HandleClient* and *HandleRef*. Procedure *HandleClient* sets the variable parameter *continue* to *false* in order to stop the search for external clients (during module unloading, one only needs to know *whether* clients exist, but not how many). Procedure *HandleRef* checks whether the argument supplied by procedures *Kernel.Scan* (either a type tag or a procedure variable within a heap record) and *FindStaticRefs* (a global procedure variable) references *any* of the selected modules to be unloaded.

We emphasize that procedure *Modules.Check* is not only called when a module is unloaded by the *user*, but also by the Oberon *background task* that removes no longer referenced *hidden* module data from memory. Thus, it must be written such that it can correctly handle *both* visible *and* hidden modules in the module data structure rooted in module *Modules*.

---

[13] *A simplified version scanning only r e c o r d blocks (allocated via NEW(p), where p is a POINTER TO RECORD) is shown. The full implementation also covers a r r a y blocks in the heap.*

Implementation prerequisites

In order to make the outlined validation pass possible, type descriptors of *dynamic* objects[14] allocated in the *heap* as well as the descriptors of *global* module data located in *static* module blocks contain a list of *procedure variable offsets*, adopting an approach employed in one of the earlier implementations of the Original Oberon system (MacOberon)[15]. The resulting run-time representation of a *dynamic* record and its associated type descriptor is shown in Figure 2.
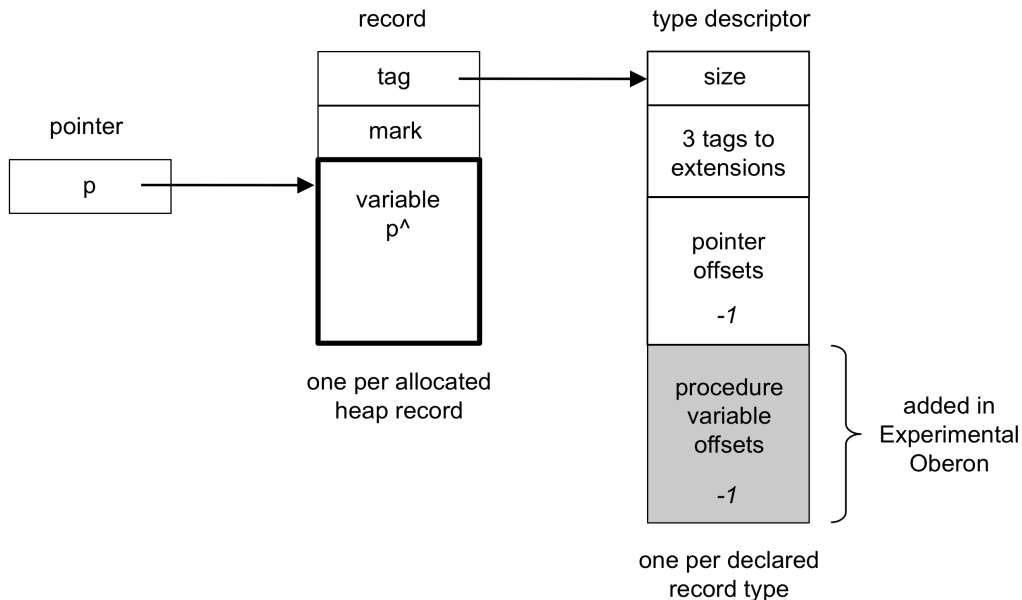
Figure 2: Run-time representation of a dynamic record and its type descriptor in Experimental Oberon

Note that this run-time layout requires the *scan* phase of reference checking to traverse over the list of *pointer offsets* in the type descriptor of *each* marked heap record. To avoid this, one could therefore use a slightly different layout, where *procedure variable offsets* are *prepended* (rather than appended) to the fields of each type descriptor[16]. We have refrained from introducing this (minor) improvement, as the resulting runtime representation might come into conflict with some implementations of the Oberon-2 programming language, which typically *also* prepend additional run-time information to the fields in each type descriptor (namely a "method table" associated with an Oberon-2 type descriptor, which however serves a different purpose than the list of procedure variable offsets used for reference checking in Experimental Oberon)[17]. A comparison of the two variants showed that the overhead in our implementation is negligible.

The descriptors also contain the offsets of *hidden* (not exported) procedure variables, enabling the module *unload* command to check *all* possible procedure variable locations in the entire system for possible procedure variable references to the modules to be unloaded.

To make the offsets of hidden *procedure variables* in exported record types available to client modules, symbol files also include them. An importing module may, for example, declare a global variable of an imported record type, which contains *hidden* procedure variables, or declare a record type that contains or extends an imported one. We recall that in Original

---

[14] See chapter 8.2, page 109, of the book Project Oberon 2013 Edition for a detailed description of an Oberon type descriptor. A type descriptor contains certain information about dynamically allocated records that is s h a r e d by all allocated objects of the same record type (such as its size, information about type extensions and the offsets of pointer fields within the record).
[15] http://e-collection.library.ethz.ch/eserv/eth:3269/eth-3269-01.pdf (The Implementation of MacOberon, 1990)
[16] Implementing this would require only a one-line change to both the compiler (procedure ORG.BuildTD) and the scan phase of reference checking (procedure Kernel.Scan).
[17] One could also adapt the Oberon-2 implementation, such that the method table is allocated somewhere else (e.g., in the heap at module load time).

Oberon, hidden *pointers*, although not exported and therefore invisible in client modules, are already included in symbol files because their offsets are needed for garbage collection[18]. Similarly, in Experimental Oberon, the locations of visible *and* hidden procedure variables are needed for reference checking during module unloading, as shown in the following example.

```
1   MODULE M0;
2     TYPE Proc* = PROCEDURE; Rec* = RECORD p*, q: Proc END ;   (*q is a hidden field*)
3     PROCEDURE Q*(r: Rec); BEGIN r.q END Q;
4     PROCEDURE Init*(VAR r: Rec; p, q: Proc); BEGIN r.p := p; r.q := q END Init;
5   END M0.
6
7   MODULE M1;
8     IMPORT M0;
9     VAR r: M0.Rec;
10    PROCEDURE Init*(p, q: M0.Proc); BEGIN M0.Init(r, p, q) END Init;
11  END M1.
12
13  MODULE M2;
14    IMPORT M1;
15    PROCEDURE Q; BEGIN END Q;
16    PROCEDURE Set1*; BEGIN M1.Init(NIL, NIL) END Set1;
17    PROCEDURE Set2*; BEGIN M1.Init(NIL, Q) END Set2;
18  END M2.
19
20  M2.Set1  System.Free M2 ~        … can unload M2
21  M2.Set2  System.Free M2 ~        … can't unload M2 (as M2.Q is referenced in global procedure variable M1.r.q)
```

Here the global record *r* declared in module *M1* contains a hidden field *M1.r.q* that is accessible only in the exporting module *M0*. If another module *M2* installs a procedure *M2.Q* in this hidden record field, for example by using an installation procedure *Init* provided by *M0* and called in *M1*, a procedure variable reference from *M1* to *M2* is created, with the (intended) effect that module *M2* can no longer be unloaded.

Hidden *pointers* are included in symbol files without their names, and their (imported) type in the symbol table is of the form *ORB.NilTyp*. Hidden *procedure variables* are also included without their names, and their symbol table type is of the form *ORB.NoTyp*[19]. For additional details, see procedures *ORB.FindHiddenFields*, *ORG.FindRefFlds*, *ORG.NofRefs* and *ORG.FindRefs*.

An obvious shortcoming of the reference checking scheme outlined above is that it requires *additional* run-time information to be present in *all* type descriptors of *all* modules *solely* for the purpose of reference checking, which in turn is *only* needed in the relatively infrequent case of module releases. In addition, modules now also contain an *additional* section in the module block containing the offsets of global procedure variables. However, since there exists only one type descriptor per declared record *type* rather than one per allocated heap *record*, and global procedure variables tend to be rare, the additional memory requirements are negligible[20].

---

[18] See chapter 12.6.2, page 41, of the book Project Oberon 2013 Edition for a description of why the offsets of hidden pointers are needed during garbage collection.
[19] This is acceptable because for record fields, the types ORB.NilTyp and ORB.NoTyp are not used otherwise.
[20] Adding procedure variable offsets to type descriptors is, strictly speaking, not necessary, as the compiler could always rearrange the memory layout of a record such that all procedure variable fields are placed in a contiguous section at the beginning of it, making their offsets implicit. Thus, it is in principle possible to realize the reference checking scheme without additional fields in type descriptors. We refrained from implementing this refinement for two main reasons. First, it would increase the complexity of the compiler. At first glance, it might seem that in order to rearrange a record's field list, only a small change needs to be made to a single procedure in the compiler (ORP.RecordType in Oberon 2013). However, a complication arises because records may recursively contain other records, each again containing procedure variables at arbitrary offsets. While it is always possible to "flatten out" such recursive record structures, it would make other record operations more complex. For example, assignments to subrecords would become less natural, as their fields would no longer be placed in a contiguous section in memory. Second, the memory savings in type descriptors would be marginal, given that there exists only one type descriptor per declared record t y p e rather than one per allocated heap r e c o r d. Most applications are programmed in the conventional programming style, where installed procedures are rare. For example, in the Oberon operating system, the object-oriented programming style is restricted to the viewer system, which provides distributed control in the form of installed handlers – of which there is usually one per t y p e of viewer. In sum, the benefit obtained by saving a few fields in a relatively small number of type descriptors appears negligible, and the additional effort required to implement this refinement would be hard to justify.

Although the operation of *reference checking* may appear expensive at first, in practice there is no performance issue. It is comparable in complexity to garbage collection and thus barely noticeable – at least on systems with small to medium sized dynamic spaces. This is in spite of the fact that for *each* heap record encountered in the *mark* phase *all* modules to be unloaded are checked for references during the *scan* phase. Note that reference checking *stops* when the *first* reference is detected, and module unloading is rare except, perhaps, during development, where however the *number* of references tends to be small.

Thus, the presented solution, which was mainly chosen for its simplicity, appears to be amply sufficient for most practical purposes.

Alternatives considered

As an alternative to the outlined conventional mark-scan scheme – where the *mark* phase first marks *all* heap objects reachable by *all other* loaded modules and the *scan* phase then checks for references from the *marked* records to the modules to be unloaded – one might be tempted to check for references directly *during* the *mark* phase and simply *stop* marking records as soon as the first reference is found. While this has the potential to be more efficient, it would lead to a number of complications.
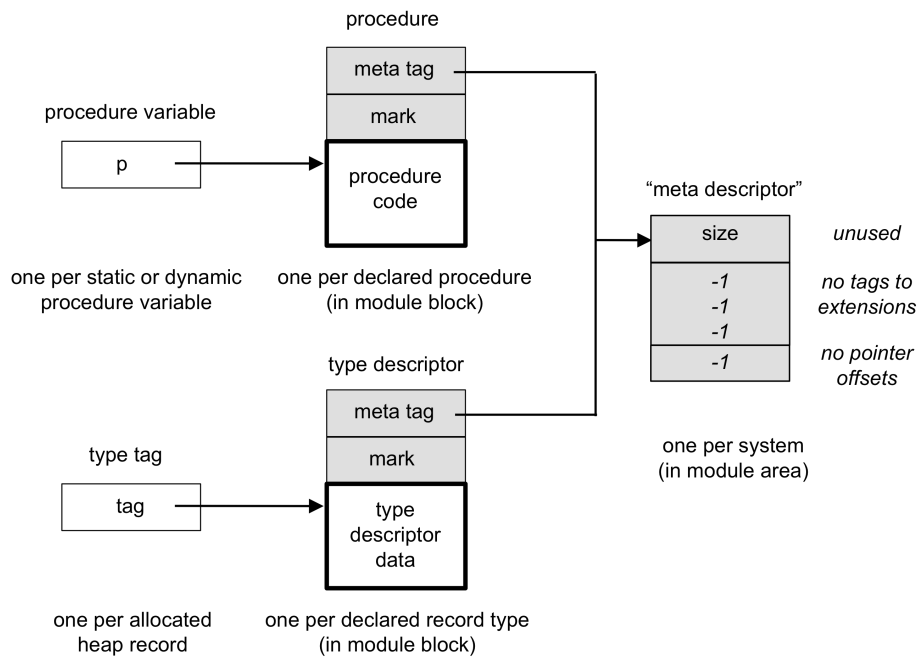
First, we note that the pointer rotation scheme used during the *mark* phase temporarily modifies not only the *mark* field, but also the *pointer variable* fields of the encountered heap records (the Deutsch-Schorr-Waite graph marking algorithm[21] is used, which essentially establishes a "return path" used during future visits of a node, effectively replacing the stack of procedure activations by a stack of marked nodes). As a consequence, one cannot simply exit the *mark* phase when a reference is found, but would also need to undo *all* pointer modifications made up to that point. The easiest way to achieve this is by letting the *mark* phase run to completion. But this would undo the performance gain initially hoped for. Second, if one wants to omit in module *Kernel* any reference to the data structure managed by module *Modules*, one would also need to express the *mark* phase as a *generic* heap traversal scheme accepting parametric handler procedures for reference checking, similar to the generic heap *scan* scheme outlined above. But such a generalization would open up the possibility for an erroneous handler procedure to prematurely end the *mark* phase, leaving the heap in a potentially irreparable state. In sum, one seems well advised not to meddle with the *mark* phase. Finally, one would still need a separate *scan* phase to *unmark* the already marked portion of the heap.

Another possible variant would be to treat *procedure variables* and *type tags* like *pointers*, and *procedures* and *type descriptors* referenced by them like *records* during the *mark* phase of reference checking. This could easily be achieved by making procedures and type descriptors *look like* records, which in turn could be accomplished by making them carry a *type tag* and a *mark field*. These additional fields would be inserted as a prefix to (the code section of) *each* declared procedure and (the type descriptor of) *each* declared record type in the module block. All static *procedures* would share a common "procedure descriptor" and all *type descriptors* a common "meta-type descriptor". These descriptors would contain no tags to extensions and no pointer offsets. Consequently, they can be represented by one and the same shared global "meta descriptor", which could be stored at a fixed location within the module area for example. The resulting run-time representation of *procedures* and *type descriptors* is shown in Figure 3.
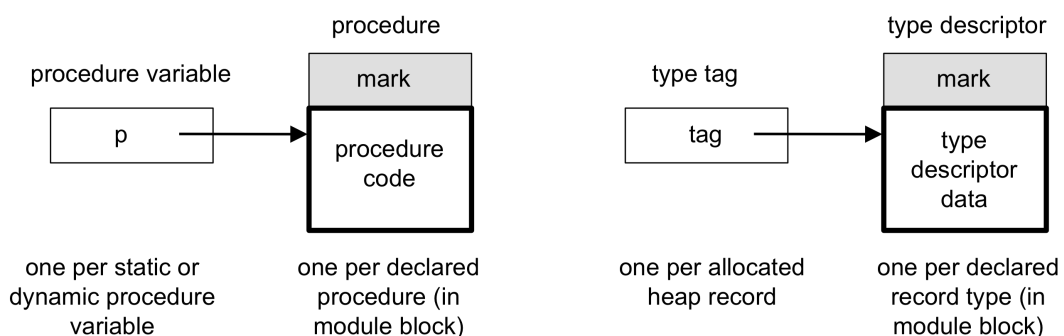
[21] Herbert Schorr and William M. Waite, "An efficient machine-independent procedure for garbage collection in various list structures", CACM, 10(8):501-506, August 1967.

Figure 3 (diagram)

procedure

meta tag
mark
procedure
code

procedure variable

p

one per static or dynamic
procedure variable

one per declared procedure
(in module block)

"meta descriptor"

size — unused
-1 — no tags to
-1 — extensions
-1
-1 — no pointer offsets

one per system
(in module area)

type descriptor

meta tag
mark
type
descriptor
data

type tag

tag

one per allocated
heap record

one per declared record type
(in module block)

Figure 3: Procedure variable and type tag interpreted as *pointer*, and procedure and type descriptor interpreted as *record*

A simpler variant would be to treat procedure variables and type tags as *special cases* during the *mark* phase, eliminating the need for a *meta* tag field as well as the shared *meta* descriptor. The resulting run-time representation of *procedures* and *type descriptors* is shown in Figure 4.

Figure 4 (diagram)

procedure

mark
procedure
code

procedure variable

p

one per static or
dynamic procedure
variable

one per declared
procedure (in
module block)

type descriptor

mark
type
descriptor
data

type tag

tag

one per allocated
heap record

one per declared
record type (in
module block)

Figure 4: Procedures and type descriptors with an additional *mark* field

With these preparations, the *mark* phase of reference checking could be suitably *extended* to also include *procedure variables* in the list of "pointers" to be traversed. In sum, the mark phase not only "touches" all reachable objects in the *heap*, but also those in the static *module blocks*.

While this technique appears appealing at first glance, a few points are worth mentioning. First, the *extended* mark phase requires an extra *mark* field to be inserted as a prefix to *each* procedure and *each* type descriptor in the module block. Given that most procedures and type descriptors are never referenced, this appears to be overkill. One could therefore decide to add the *mark* field only to *module descriptors* rather than individual *procedures* or *type descriptors* during the *mark* phase. While this would make the check whether any of the selected modules is referenced a trivial task, it would render the *mark* phase more complex, as it would now need to *locate* the module descriptor belonging to a given procedure or type descriptor (on systems where additional meta-information is present in the run-time representation of loaded modules, such as the locations of procedures and type descriptors in module blocks or *back pointers* from

each object of a module to its module descriptor, the mark phase would be simpler; however, neither Original Oberon 2013 nor Experimental Oberon offer such *metaprogramming* facilities). Second, one would still need to mark *all* objects, for the same reason as outlined above, i.e. one cannot simply exit in the middle of the *mark* phase without additional action. Third, one would also still need an extra *scan* phase to unmark the marked objects.

Finally, a comparison of the code required to implement the various alternatives showed that our solution is *by far* the simplest: the combined implementation cost of *all* modifications to the runtime representation of type descriptors and descriptors of global module data, the object and symbol file formats and the module loader, is only about 15 source lines of code[22], while the *reference checking* phase itself amounts to less than 75 lines (the *total* implementation cost to realize *safe module unloading*, including the ability to unload module *groups* and the automatic collection of no longer referenced hidden modules, amounts to about 250 source lines of code).

* * *

---

[22] *See procedures ORB.InType (+ 1 line), ORB.FindHiddenFields (+ 1 line), ORG.BuildTD (+ 1 line), ORG.Close (+ 7 lines) and Modules.Load (+ 4 lines).*

Appendix: Historical notes on module unloading in the Oberon system

In most Oberon implementations, only *client* references are checked prior to module unloading, i.e. if clients exist among the remaining modules, a module or module group is not unloaded.

As for *type* and *procedure variable* references, various approaches have been used to address them. For such references, there exist two essentially different interpretations of the semantics of module unloading:

a. Schemes that explicitly allow invalidating past references

In such schemes, removing a module from memory *may*, and in general *will*, lead to "dangling" references, i.e. references that point to module data that is no longer valid. Such references can be in the form of *type tags* (addresses of type descriptors) in *dynamic* objects or in the form of *procedure variables* installed in *static* or *dynamic* objects[23].

An important use case is when a structure rooted in a variable of base type T declared in a base module M (for example module *Viewers*) contains elements of an extension T' defined in a client module M' (for example a graphics editor), which is then unloaded. Such elements typically contain both *type* references (type tags) and *procedure variable* references (installed handler procedures) that still refer to M'.

In addition, *global* procedure variables declared in other modules may also refer to procedures in module M', although this case is less common (global procedure variables tend to be used mainly for procedures declared in the *same* module).

A variety of approaches have been used in various implementations of the Oberon system to cope with the introduced dangling *type* or *procedure variable* references:

1. The easiest way to cope with dangling references is to simply *ignore* them. This is the approach chosen in FPGA Oberon on RISC, where the memory associated with a module to be unloaded is *always* released (unless clients exist) without taking any further precautions. But this leaves the system in an *unsafe* state. It will become *unstable* the very moment another module loaded later *overwrites* the previously released module block *and* other loaded modules still refer to its *type descriptors* or *procedures*. This is of course undesirable.

2. Another approach, which can however be used only for *procedure variable* references, is to identify *all* procedure variable references to the module M to be unloaded and make them refer to a "dummy" procedure, preventing a run-time error, when such "fixed up" procedures are called later. This solution was tested in an earlier version of Experimental Oberon, but was later discarded, mainly because the resulting effect on the *overall* behavior of the system would be essentially impossible to predict (or even detect) by the end user. The fact that *some* procedure variables *somewhere* in the system no longer refer to "real", but to "dummy" procedures typically becomes "visible" only through the *absence* of some action – such as mouse tracking if the unloaded module contained a viewer handler, for example.

3. On systems that use *indirection* for procedure calls via a so-called "link table"[24], one can achieve the same effect by setting the *link table entries* for all referenced procedures of a

---

[23] *If the programming language Oberon-2 is used, there can also be references to method tables (which however are typically allocated within type descriptors themselves).*
[24] *When a system uses i n d i r e c t i o n via a link table, an "address" of a procedure is not a real memory address, but an index to this translation table – which the caller consults for every procedure call, in order to obtain the actual memory location of the called procedure.*

module to be unloaded to *dummy* entries, instead of locating and modifying each individual procedure *call* anywhere in the system.

We note that using a link table to implement indirection for procedure calls is only viable on systems that provide *efficient* hardware support for it. It has been used in some of the earlier versions of the Oberon system on Ceres computers, which were based on the (now defunct) NS32000 processor. Early versions of this processor featured efficient *hardware* support for the link table in the form of a *call external procedure* (CXP k) instruction (where k is the index of the link table entry of the called procedure), which sped up the process of activating external procedures significantly[25]. Later versions of the NS processor, however, internally re-implemented the same processor instruction using microcode. This negatively impacted its performance – so much that it became slower than the *regular branch to subroutine* (BSR) instruction[26]. Therefore, the CXP k instruction – and with it the *link table* – were no longer used in later versions of Ceres-Oberon, for example on Ceres-3.

4. On systems that use a *memory management unit* to perform virtual memory management, such as on Ceres-1 or Ceres-2, another possible approach is to simply *unmap* the module space of an unloaded module from virtual memory, thereby *invalidating* future references to it. *After* that, a dangling reference points to a now unallocated page, and consequently any attempt to access this page, for example via *type* or *procedure variable* references, results in a *trap* on Ceres-1 and Ceres-2, thereby preventing a system crash.

We consider this an unfortunate proposal for several reasons. First, users generally have no way of knowing *whether* it is in fact safe to unload a module, yet they are allowed to do so. Second, after having unloaded it, they still don't know whether references from other loaded modules have existed or still exist – until a *trap* occurs. Only then they know.

Another disadvantage of this solution is that it requires special hardware support, which may not be available on all systems. Indeed, on Ceres-3, which does not use virtual memory, an attempt to access an unloaded module M goes undetected *initially* – until a *trap* or, worse, a *crash* occurs later. This can, and usually *will*, happen, the moment another module is loaded into the module block previously occupied by the unloaded module M *and* the overwritten data is still *referenced* by other modules[27].

5. Finally, we add the remark that for *type* references, it is actually possible to determine at *compile* time, whether a module M will *not* lead to references from other loaded modules at run time: namely, if module M does *not* declare record types which are extensions T' of an imported type T, then records declared in M *cannot* ever be inserted in a data structure rooted in a variable *v* of an imported type T – precisely *because* they are not extensions of T (in the Oberon programming language, the assignment *p := p'* is allowed only if the type of p' is the same as the type of *p* or an extension of it).

One *could* therefore introduce a rule that a module M can be safely dispensed *only* if it does *not* declare record types, which are extensions T' of an imported type T. The flip side of such a rule, however, is that modules that actually *do* declare such types can *never* be unloaded.

---

[25] *The use of the link table also increased code density considerably (as only 8 bits for the index instead of 32 bits for the full address were needed to address a procedure in every procedure call). In addition, the link table used by the CXP instruction allowed for an expedient linking process at load time (as there are far fewer conversions to be performed by the module loader – one for every referenced procedure instead of one for every procedure call) and also eliminated the need for a fixup list (list of the locations of all external procedure references to be fixed up by the module loader) in the object file. A disadvantage is, of course, the need for a (short) link table.*

[26] *The internal re-implementation of the CXP instruction using microcode in later versions of the NS processor followed the general industry trend of implementing only frequent, simple instructions directly with hardware, while interpreting more complex instructions using internal microcode. In general, with the advent of highly regular reduced instruction set computers (RISCs) in the 1980s and 1990s, the trend towards offering microprocessors providing a smaller set of simple instructions, most of them executing in a single clock cycle, combined with fairly large banks of (fast) registers, continued – and does so to this date.*

[27] *Of course, one could adapt the module loader to never overwrite a previously "unloaded" module block, even if a new module to be loaded would fit, but this would be a waste of memory .*

Even though most of these approaches have actually been realized in various implementations of the Oberon system, we consider none of them truly satisfactory. In our view, these schemes appear to only tinker with the symptoms of a problem that would not exist, if only one adopted the rule to *disallow* the removal (from memory) of still referenced module data.

The main issue is that the moment one *allows* modules to release their associated memory if references to them still exist, the resulting *dangling* references must be "fixed up" *somehow*, in order to prevent an almost certain system *crash* (when their module block is overwritten).

However, fixing up references will *always* remove *essential* information from the system. As a result, the run-time behavior of the modified system becomes *essentially unpredictable*, as other loaded modules may *critically* depend on the removed functionality.

For example, unloading a module that contains a handler procedure of a *contents frame* may render it impossible to *close* the enclosing *menu viewer* that contains it. If the approach chosen to handle dangling procedure variable references involves generating a trap[28], the trap itself will actually *prevent* a system crash (as intended), but the user may *still* need to reboot the system in order to recover an environment without any "frozen" parts, e.g. still open viewers[29].

A similar problem may occur with references to *type descriptors*, if they are not preserved in memory *after* unloading their associated modules.

b. Schemes where past references remain unaffected

The second possible interpretation of *unloading* a module consists of schemes where *past* references *must* at all times remain *unaffected*. In such schemes, module unloading can be viewed as an implicit mandate to preserve "critical module data", as long as references exist.

1. One could of course simply exit the *unload* command with an error message, whenever such references are detected. The user, however, may then be "stuck" with modules that he can *never* unload because they are referenced by modules over which he has no explicit control.

2. But the mandate could also be fulfilled by allowing the user to *persist* any still referenced module data to a "safe" location before unloading the associated module.

    For *type* references (type tags referring to type descriptors) a straightforward solution exists: simply allocate type descriptors *outside* the module blocks, in order to persist them beyond the lifetime of their associated module. One possibility is to allocate them in the *heap* at module load time[30]. This has been implemented in Oberon on Ceres-3, for example. Note that this method eliminates dangling *type* references altogether and therefore also the *need* to check for them at run time.

    For *procedure variable* references no such simple solution exists. The only way to persist procedures would be to persist the *entire* module (recall that procedures may *access* global module data or *call* other procedures of the same module).

---

[28] *In that case, if the enclosing menu viewer attempts to send a "close" message to the sub frame by calling its handler, a trap is generated, if the module has been unloaded before.*
[29] *The module could of course provide a "close" command, which also accepts the marked viewer as argument (using procedure Oberon.MarkedViewer), but that is not necessarily the case.*
[30] *Note that one cannot simply move type descriptors around in memory, as their addresses are (typically) used to implement type tests and type guards. By allocating them in the heap at module load time, one avoids the need to move them to a different location when a module is unloaded.*

We conclude that *if* one wants to address both type *and* procedure variable references, one *cannot* unload the module block from memory, as long as references to it still exist[31].

3. One possible approach consists of simply *not* releasing the associated memory of a module to be unloaded, but removing it only from the *list* of currently loaded modules. This amounts to *renaming* the module[32], with the implication that a newer version of the same module with the same name can be reloaded again. Such an approach has been implemented in MacOberon[33], for example. Since the associated memory of a module is *never* released, the issue of dangling type or procedure variable references is avoided altogether, as they simply cannot exist. However, it can also lead to higher-than-necessary memory usage, if a module is repeatedly loaded and unloaded (which is typical during *development*).

    Nevertheless, such an approach may be viewed as adequate on *production* systems or on systems that use *virtual memory* with *demand paging*. On such systems, the virtual address space is (practically) unlimited and data that is accessed only infrequently is temporarily *swapped out* from main to secondary store. Thus, once a module has been removed from the module list, its module block is – over time – less likely to be accessed and hence more likely to be placed in secondary store by the demand paging mechanism, freeing up valuable memory resources for modules loaded later. Of course, if the *swapped out* module is accessed later by another module that still references it, it will have to be swapped back in.

    Note that with the advent of large primary stores, the concept of virtual memory with demand paging has lost much of its significance. Therefore it is not necessarily used on all systems. On systems that do *not* use virtual memory (such as Ceres-3 or FPGA Oberon on RISC), other mechanisms to handle no longer referenced module data must be considered, if one wants to avoid running out of memory *eventually*.

4. A *refinement* of the approach outlined above consists of *initially* removing a module from the *list* of loaded modules (as is done in MacOberon), but *in addition* releasing its associated memory *as soon as* there are no more references to it. If this is done in an automatic fashion (for example as part of a background process), module data is truly "kept in memory for exactly as long as necessary and removed from it as soon as possible".

    This is the approach chosen in Experimental Oberon. A variation of it was used in one of the later versions of SparcOberon[34]. Of course, this scheme also works on systems that *do* use a memory management unit to perform virtual memory management (it optimizes them).

In sum, schemes were past references remain *unaffected* avoid many of the complications that are inherent in schemes that explicitly *allow* invalidating past references. A (small) price to pay is to keep loaded modules in memory, as long as references to them exist.

However, on *production* systems, there is typically *no need* to keep multiple copies of the same module loaded in memory, while on *development* systems it is *totally acceptable*. Finally, we note that on modern computers, the amount of available memory – and consequently the amount of *dynamic* data allocated by modules – typically far exceeds the size of their *static* module blocks. Hence, not releasing module blocks immediately has a rather negligible impact.

---

[31] Of course, a "mixed" variant is also possible, namely to allocate type descriptors in the heap, and preserve a module block o n l y if procedure variable references exist; however, most modules referenced by type tags are also referenced by procedures – this is in fact the typical case for records with installed handlers. Thus, this would be an "optimization" in the wrong place.
[32] In a specific implementation, one might choose to make the module completely anonymous or modify the name such that one can no longer import it (e.g., by inserting an asterisk).
[33] http://e-collection.library.ethz.ch/eserv/eth:3269/eth-3269-01.pdf (The Implementation of MacOberon, 1990)
[34] http://e-collection.library.ethz.ch/eserv/eth:7103/eth-7103-01.pdf (SPARC-Oberon User's Guide and Implementation, 1990/1991)