

Safe module unloading

in the Oberon operating system

Andreas Pirklbauer

15.7.2019

Purpose

This technical note describes an implementation of “safe” module unloading for the Oberon operating system, as realized in Experimental Oberon¹, a revision of FPGA Oberon².

Module unloading in the Oberon system

In the Oberon system, there exist three possible types of references to a loaded module M³:

1. *Client references* exist when other loaded modules import module M. Client modules may refer (by name) to constants, types, variables or procedures declared in module M.
2. *Type references* exist when *type tags* (addresses of type descriptors) in dynamic objects reachable by other loaded modules refer to descriptors of types declared in module M.
3. *Procedure variable references* exist when procedure variables in static or dynamic objects reachable by other loaded modules refer to procedures declared in module M.

In most Oberon implementations, only *client* references are checked prior to module unloading. *Type* and *procedure variable* references are usually not checked, although various approaches are typically employed to address the case where such references exist⁴. As a result, module unloading in Oberon has traditionally left the system in an *unsafe* state, which will become *unstable* the moment a module loaded later *overwrites* a previously released module block *and* other loaded modules still refer to its *type descriptors* or *procedures*. In addition, unloading of module *groups*, which contain references only among themselves, is usually not supported.

¹ <http://www.github.com/andreaspirklbauer/Oberon-experimental>

² <http://www.projectoberon.com>

³ An Oberon module can be viewed as a container of constants, types, variables and procedures, where variables can be procedure-typed. Global objects with an explicit name declared in a module are allocated in the module area when the module is loaded. If exported, such objects may be referenced by name in client modules. Anonymous variables with no explicit name declared in a module are allocated in the heap when needed during program execution using the predefined procedure NEW. Such objects may be reachable by multiple pointer variables, potentially declared in different modules. Pointer types are bound to their base types, leading to additional hidden type references, thereby making type checks on dynamic objects possible. Assignments to procedure-typed variables lead to procedure variable references via installed procedures. In sum, there can be type, variable, pointer, procedure and procedure variable references from static or dynamic objects of other modules to static or dynamic objects of the modules to be unloaded. However, only client references, hidden type references from dynamic objects to types declared in the modules to be unloaded and procedure variable references from static or dynamic objects to procedures declared in the modules to be unloaded need to be checked prior to module unloading for the following reasons. First, exported constants, types, variables and procedures referenced by name in client modules are already handled via their import/export relationship (if clients exist, a module or module group is never unloaded) and therefore don't need to be checked separately. Second, pointer references from static or dynamic pointer variables reachable by other modules to dynamic objects reachable by the modules to be unloaded should not prevent module unloading and therefore should not be checked. Such references will be handled by the Oberon garbage collector during a future garbage collection cycle (dynamic objects that were reachable only by the unloaded modules prior to their unloading will be collected). Finally, pointer references to static objects of the modules to be unloaded are only possible by resorting to low-level facilities and should be avoided – and, in fact, be disallowed (pointers should point exclusively to anonymous variables dynamically allocated in the heap when needed during program execution).

⁴ See the appendix for historical notes on module unloading in the Oberon system.

Implementing “safe” module unloading

In Experimental Oberon, a revision of FPGA Oberon⁵, *all* possible types of references to a loaded module or module group are checked as follows prior to module *unloading* (Figure 1):

- If clients exist among the loaded modules, a module or module group is never unloaded.
- If no client, type or procedure variable references to a module or module group exist in the remaining modules or data structures, it is unloaded and its associated memory is released.
- If no clients, but type or procedure variable references exist, the module *unload* command takes no action by default and merely displays the names of the modules and the type of references that caused the removal to fail. If, however, the *force* option */f* is specified⁶, the modules are initially removed only from the *list* of loaded modules, without releasing their associated memory. Such “hidden” modules are later physically removed from *memory* using the command *Modules.Collect*⁷, as soon as there are no more references to them.

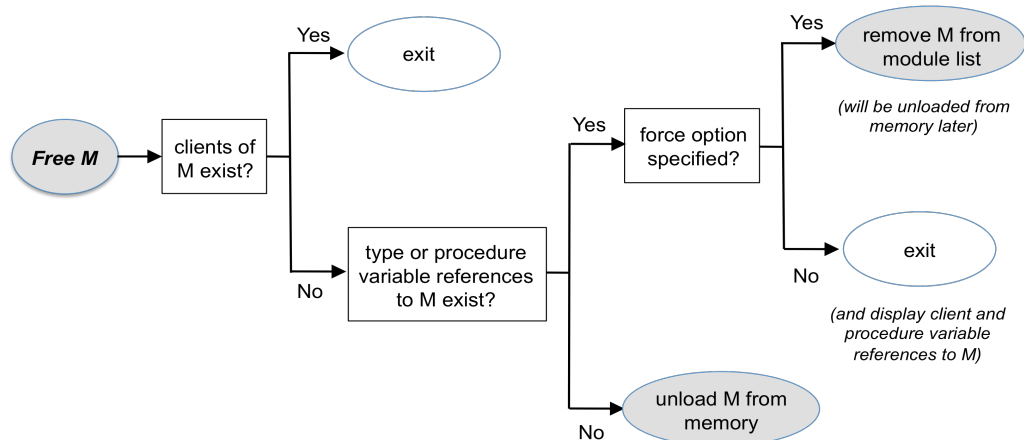


Figure 1: Safe module unloading in Experimental Oberon

Removing a module only from the module *list* amounts to *renaming* it, with the implication that a newer version of the same module (with the same name) can be reloaded again, without having to unload (from memory) earlier versions that are still referenced⁸. Unloading a module from *memory* frees up the memory area previously occupied by the module block⁹. In Experimental Oberon, this is only possible when no references to the module exist.

To automate the process of unloading no longer referenced *hidden* module data, the command *Modules.Collect* has been included in the Oberon background task handling garbage collection. It checks all possible combinations of *k* modules chosen from *n* hidden modules for references to them, and removes those module subgroups from memory that are no longer referenced.

In sum, module unloading does not affect any references, as module data is kept in memory for exactly as long as necessary and removed from it as soon as possible.

⁵ <http://www.projectoberon.com>

⁶ The force option */f* must be specified at the *e n d* of the list of modules to be unloaded, e.g., *System.Free M1 M2 M3/f*

⁷ The tool command *System.Collect* also invokes *Modules.Collect*.

⁸ Modules removed only from the list of loaded modules are marked with an asterisk in the output of the command *System.ShowModules*. Commands of such “hidden” modules can be accessed by either specifying their module number or their (modified) module name, both of which are displayed by the command *System.ShowModules*. In both cases, the corresponding command text must be enclosed in double quotes. If a module *M* carries module number 14, for instance, one can activate a command *M.P* also by clicking on the text “14.P”. Typical use cases include hidden modules that still have background tasks installed which can only be removed by activating a command of the hidden modules themselves, or hidden modules that still have open viewers. If the command to close a viewer is displayed in the viewer’s menu bar, the user can manually edit the command text (by clicking within its bottom two pixel lines), replace the module name by its module number and enclose the modified command text in double quotes. Although this is somewhat clumsy, it at least enables the user to close the viewer. An alternative approach is to provide a “Close” command that also accepts the marked viewer as argument (using procedure *Oberon.MarkedViewer*).

⁹ In FPGA Oberon and in Experimental Oberon, the module block includes the module’s type descriptors. In some other Oberon implementations, such as Ceres-Oberon, type descriptors are not stored in the module block, but are allocated in the heap at module load time in order to persist them beyond the lifetime of their associated modules. In Experimental Oberon, no such extra precaution is necessary, as module blocks are removed only from the list of loaded modules, if they are still referenced. Thus, type descriptors can safely be stored in the module blocks.

For example, older versions of a module's code can still be executed if they are referenced by static or dynamic procedure variables in other modules, even if a newer version of the same module has been reloaded in the meantime¹⁰. Type descriptors also remain accessible to other modules for exactly as long as needed. This covers the important case where a structure rooted in a variable of base type *T* declared in a base module *M* contains elements of an extension *T'* defined in a client module *M'*, which is unloaded. Such elements typically contain both *type* references (type tags) and *procedure variable* references (handler procedures) referring to *M'*.

If a module *group* is to be unloaded and there exist references *only* within this group, it is unloaded as a *whole*. This can be used to remove module groups with *cyclic* references¹¹. It is also possible to release *entire subsystems* of modules. The command *System.FreeImports* attempts to unload the specified modules and all their direct and indirect *imports*. It may be used for conventional module stacks with a single top module, for example a compiler. By contrast, the command *System.FreeClients* unloads the specified modules and all their direct and indirect *clients*. This variant may be used for module stacks written in the object-oriented programming style with their typically inverted module hierarchy, for example a graphics editor.

Note that these unloading strategies amount to *heuristics* and may tend to unload rather large parts of the system, which may not be desired. The recommended way to unload modules is to use the base command *System.Free* with a *specific* set of modules provided as parameters. For added convenience, the tool commands *System.ShowRefs* and *System.ShowGroupRefs* can be used to identify all modules containing references to a given module or module group.

Implementation aspects

To check for references from other loaded modules, the module *unload* command must first *select* the modules to be unloaded using procedure *Modules.Select* and then invoke procedure *Modules.Check*. Client references are checked first. This is accomplished by simply verifying whether *unselected* modules import *selected* modules¹²:

```

1  PROCEDURE FindClients*(proc: ImpHandler; VAR res: INTEGER);
2  VAR mod, imp, m: Module; p, q: INTEGER; continue: BOOLEAN;
3  BEGIN res := noref; m := root; continue := proc # NIL;
4  WHILE continue & (m # NIL) DO
5    IF (m.name[0] # 0X) & m.selected & (m.refcnt > 0) THEN mod := root;
6    WHILE continue & (mod # NIL) DO
7      IF (mod.name[0] # 0X) & ~mod.selected THEN p := mod.imp; q := mod.cmd;
8      WHILE p < q DO imp := Mem[p];
9      IF imp = m THEN INC(res, proc(mod, imp, continue)); p := q ELSE INC(p, 4) END
10     END
11     END ;
12     mod := mod.next
13   END
14   END ;
15   m := m.next
16   END
17 END FindClients;
```

¹⁰ If an older version of a module's code accesses global variables (of itself or of other modules), it will automatically access the "right" version of such variables.

¹¹ In Oberon, cyclic references can be created by pointers or procedure variables, whereas cyclic module imports are normally not allowed. However, through a tricky sequence of compilation and editing steps, it is in fact possible to *construct* cyclic module imports, which cannot be detected by the compiler. But even though they can be created, such modules are not allowed to be loaded, as the module loader of Experimental Oberon – adopting the approach chosen in Original Oberon and FPGA Oberon – would enter an endless recursion, eventually leading to a stack overflow or out-of-memory condition – a totally acceptable solution for such an artificially constructed case. But even if modules with cyclic module imports *were* allowed to be loaded, Experimental Oberon would handle them correctly upon unloading, i.e. if no external clients or references exist, such a module group would simply be unloaded as a whole – as it should.

¹² *Mem* stands for the entire memory and assignments involving *Mem* are expressed as *SYSTEM.GET(a, x)* for *x := Mem[a]* and *SYSTEM.PUT(a, x)* for *Mem[a] := x*.

If clients exist among the *unselected* modules, no further action is taken and the module *unload* command exits. If no clients exist, dynamic *type* and *procedure variable* references from dynamic objects in the heap¹³ are checked next using a conventional *mark-scan* scheme:

```

1  PROCEDURE FindDynamicRefs*(type, proc: RefHandler; VAR resType, resProc: INTEGER; all: BOOLEAN);
2  VAR mod: Module;
3  BEGIN mod := root;
4  WHILE mod # NIL DO
5    IF (mod.name[0] # 0X) & ~mod.selected THEN Kernel.Mark(mod.ptr);
6    IF ~all THEN Kernel.Scan(type, proc, mod.name, resType, resProc) END
7  END ;
8  mod := mod.next
9  END ;
10 IF all THEN Kernel.Scan(type, proc, "", resType, resProc) END
11 END FindDynamicRefs;

```

During the initial *mark* phase, dynamic objects reachable by all *named* global pointer variables of all *unselected* modules are marked (line 5), thereby excluding objects reachable *only* by the modules to be unloaded. This automatically recognizes module *groups* and ensures that when a module or module group is referenced *only* by itself, it can still be unloaded. The subsequent *scan phase* (line 6 or 10), implemented as a separate procedure *Scan*¹⁴ in module *Kernel*, scans the heap sequentially, unmarks all *marked* objects and checks whether the *type tags* of these objects point to descriptors of *types* and whether *procedure variables* declared in those same objects refer to *procedures* declared in the modules to be unloaded.

An additional boolean parameter *all* allows the caller to indicate whether the *mark* phase should first mark all heap objects reachable by *all* other modules before initiating the *scan* phase *once* (used for module unloading), or whether the *mark* and *scan* phases should be initiated for *each* unselected module (used for identifying references from each unselected module *individually*).

Finally, the check for *procedure variable* references is also performed for all *global* procedure variables, whose offsets in the module's data section are obtained from an array in the module's *meta* data section, headed by the link *mod.pvr* in the module descriptor (see below):

```

1  PROCEDURE FindStaticRefs*(proc: RefHandler; VAR res: INTEGER);
2  VAR mod: Module; pref, pvadr, r: LONGINT; continue: BOOLEAN;
3  BEGIN res := noref; mod := root; continue := proc # NIL;
4  WHILE continue & (mod # NIL) DO
5    IF (mod.name[0] # 0X) & ~mod.selected THEN
6      pref := mod.pvr; pvadr := Mem[pref];
7      WHILE continue & (pvadr # 0) DO r := Mem[pvadr];
8        INC(res, proc(pvadr, r, mod.name, continue));
9        INC(pref, 4); pvadr := Mem[pref]
10     END
11   END ;
12   mod := mod.next
13 END
14 END FindStaticRefs;

```

Note that the procedures *FindClients*, *FindDynamicRefs* and *FindStaticRefs* are all expressed as *generic* object traversal schemes, which accept *parametric* handler procedures that are called for each encountered object. This allows these procedures to be used for *other* purposes as well, for example to *count* or *enumerate* all references to *any* given module or module group.

¹³ In FPGA Oberon, only records can be allocated in the heap. In Experimental Oberon, fixed-length and open arrays can also be dynamically allocated.

¹⁴ The original procedure *Kernel.Scan* (implementing the scan phase of the Oberon garbage collector) has been renamed to *Kernel.Collect*, in analogy to procedure *Modules.Collect*.

In order to omit in module *Kernel* any reference to the list of modules rooted in module *Modules*, procedure *Kernel.Scan* is also expressed as a *generic* heap scan scheme. The following is a simplified version of this scheme¹⁵:

```

1  PROCEDURE Scan*(type, proc: Handler; s: ARRAY OF CHAR; VAR resType, resProc: INTEGER);
2  VAR offadr, offset, p, r, mark, tag, size: LONGINT; continue: BOOLEAN;
3  BEGIN p := heapOrg; resType := 0; resProc := 0; continue := (type # NIL) OR (proc # NIL);
4  REPEAT mark := Mem[p+4];
5  IF mark < 0 THEN (*free*) size := Mem[p]
6  ELSE (*allocated*) tag := Mem[p]; size := Mem[tag];
7  IF mark > 0 THEN (*marked*) Mem[p+4] := 0;
8  IF continue THEN
9    IF type # NIL THEN INC(resType, type(p, tag, s, continue)) END ;
10   IF continue & (proc # NIL) THEN offadr := tag + 16; offset := Mem[offadr];
11   WHILE offset # -1 DO (*skip pointers*) INC(offadr, 4); offset := Mem[offadr] END ;
12   INC(offadr, 4); offset := Mem[offadr];
13   WHILE continue & (offset # -1) DO (*procedures*) r := Mem[p+8+offset];
14   INC(resProc, proc(p+8+offset, r, s, continue));
15   INC(offadr, 4); offset := Mem[offadr]
16   END
17   END
18   END
19   END
20   END ;
21   INC(p, size)
22   UNTIL p >= heapLim
23   END Scan;
```

This scheme calls *parametric* handler procedures for individual elements of each *marked* heap object rather than *directly* checking whether they contain type or procedure variable references to the modules to be unloaded. Procedure *type* is called with the memory address of the *heap object* itself as the “source” of the reference followed by the *type tag* (address of the type descriptor) as its “destination” (line 9). Procedure *proc* is called for each procedure variable declared in the same object with the address of the *procedure variable* as the source followed by the address of the referenced *procedure* as the destination (line 14). The results of these handler calls are *separately* added up for each handler procedure and returned in the variable parameters *resType* and *resProc*.

An additional boolean variable parameter *continue* allows the handler procedures to indicate to the caller that they are no longer to be called (lines 8, 10, 13). Note that the scan process itself continues, but only to *unmark* the remaining marked heap objects (line 7).

Procedure *Modules.Check* uses procedures *FindClients*, *FindDynamicRefs* and *FindStaticRefs* by passing its own private handler procedures *HandleClient* and *HandleRef*. These procedures merely set the boolean parameter *continue*, depending on whether a reference exists or not.

We emphasize that procedure *Modules.Check* is not only called when a module is unloaded by the *user*, but also by the Oberon *background task* that removes no longer referenced *hidden* module data from memory. Thus, it must be able to correctly handle visible *and* hidden modules in the data structure rooted in module *Modules*.

¹⁵ The simplified version scans only *record* blocks allocated via *NEW(p)*, where *p* is a *POINTER TO RECORD*. The full implementation also covers *array* blocks allocated in the heap.

Implementation prerequisites

In order to make the outlined validation pass possible, type descriptors of *dynamic* objects¹⁶ allocated in the heap and the descriptors of *global* module data located in the module blocks contain a list of *procedure variable offsets*, adopting an approach employed in an earlier Oberon implementation (MacOberon)¹⁷. The resulting run-time representations of a *dynamic* record with its associated type descriptor and the modified module block are shown in Figure 2 (the *method table* used to implement Oberon-2 style *type-bound procedures* is not discussed here).

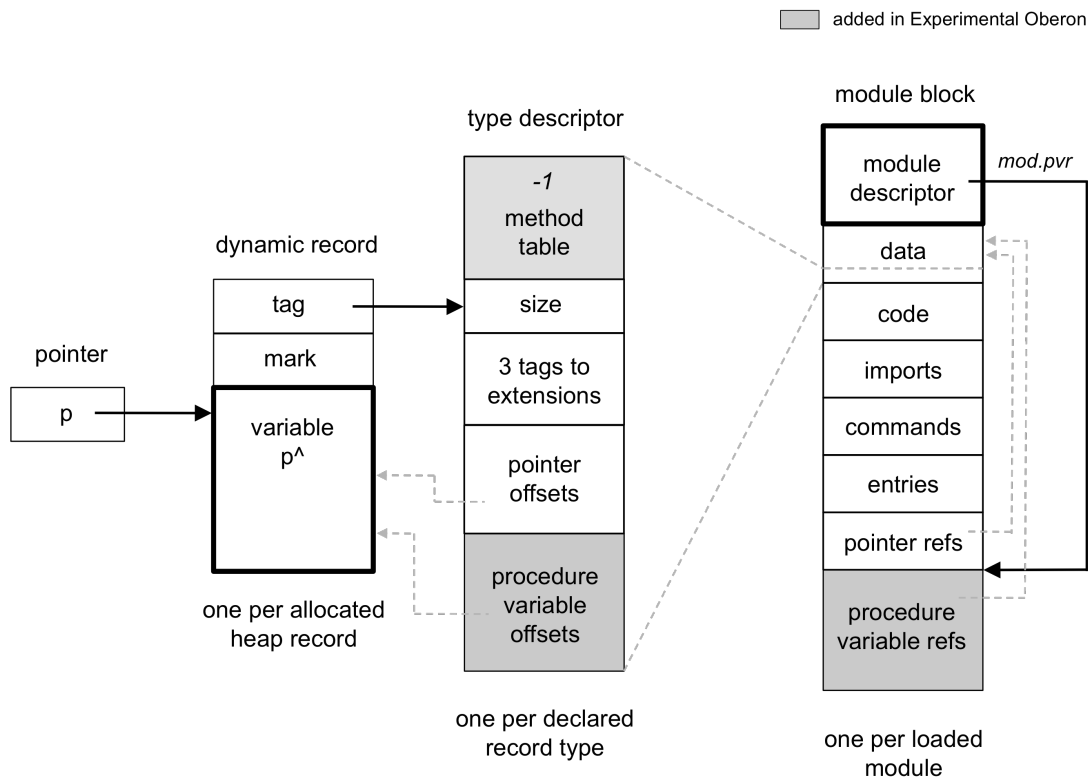


Figure 2: Run-time representation of a dynamic record with its type descriptor and a module block in Experimental Oberon

The descriptors also contain the offsets of *hidden* (not exported) procedure variables, enabling the module *unload* command to truly check *all* static and dynamic procedure variable locations in the entire system for possible procedure variable references to the modules to be unloaded.

To make the offsets of hidden *procedure variables* in exported record types available to clients, symbol files also include them. An importing module may, for example, declare a variable of an imported record type, which contains *hidden* procedure variable fields, or declare a record type, which contains or extends an imported type. We recall that in FPGA Oberon, hidden *pointers*, although not exported and therefore invisible in client modules, are nevertheless included in symbol files because their offsets are needed for garbage collection¹⁸. Similarly, in Experimental Oberon, hidden *procedure variables* are included in symbol files because their offsets are needed for reference checking during module unloading, as shown in the example below.

¹⁶ See chapter 8.2, page 109, of the book *Project Oberon 2013 Edition*. In essence, an Oberon type descriptor contains information about dynamically allocated records that is shared by all allocated objects of the same record type, such as its size, information about type extensions and the offsets of pointer and procedure variable fields within the record.

¹⁷ <http://e-collection.library.ethz.ch/eserv/eth:3269/eth-3269-01.pdf> (The Implementation of MacOberon, 1990)

¹⁸ See chapter 12.6.2 ("Symbol files"), page 41, of the book *Project Oberon 2013 Edition*, available at <http://www.projectoberon.com>.

```

1  MODULE M0;
2    TYPE P* = PROCEDURE; R* = RECORD p: P END ;           (*hidden record field p, visible only in M0*)
3    PROCEDURE Init*(VAR r: R; p: P); BEGIN r.p := p END Init; (*install procedure p in hidden field r.p*)
4  END M0.
5
6  MODULE M1;
7    IMPORT M0;
8    VAR r: M0.R;                                           (*global variable r with hidden field r.q*)
9    PROCEDURE Init*(p: M0.P); BEGIN M0.Init(r, p) END Init;
10  END M1.
11
12  MODULE M2;
13  IMPORT M1;
14  PROCEDURE P; BEGIN END P;
15  PROCEDURE ClearRef*; BEGIN M1.Init(NIL) END ClearRef;    (*clear reference from M1.r.p*)
16  PROCEDURE SetRef*; BEGIN M1.Init(P) END SetRef;          (*create reference from M1.r.p to M2.P*)
17  END M2.
18
19  M2.ClearRef System.Free M2 ~   unloading sucessful (no references to M2)
20  M2.SetRef   System.Free M2 ~   unloading failed (procedures of M2 in use in global procedure variables of M1)

```

Here the global variable *r* declared in module *M1* is of the imported record type *M0.R*, which contains a hidden procedure variable field *p*. If another module *M2* installs a procedure *M2.P* in this field (line 16), a global procedure variable reference from *M1.r.p* to *M2.P* is created, with the effect that *M2* can no longer be unloaded (line 20). In order to be able to check for such hidden references, the location of *M1.r.p* must be known at run time. This is achieved as follows:

- The offset of the hidden field *R.p* is included in the symbol file of *M0*.
- The location of the procedure variable *M1.r.p* within the data section of *M1* (computed as the sum of the starting address of *M1.r* and the offset of *R.p*) is included in the object file of *M1*.
- The location of *M1.r.p* is transferred from the object file of *M1* to the corresponding array in the *meta* data section in the module block, headed by the link *mod.pvr*, when *M1* is loaded.

Hidden *pointers* are included in symbol files without name, and their imported type in the symbol table is of the form *ORB.NilTyp* (as in FPGA Oberon). Hidden *procedure variables* are also included in symbol files without name, but their imported type in the symbol table is of the form *ORB.NoTyp*¹⁹. Details are to be looked up in the compiler procedures *ORB.FindHiddenFields*, *ORG.FindRefFlds*, *ORG.NoRefs* and *ORG.FindRefs*. The above choice for the imported type of hidden objects is reflected in the conditions used to find *all* pointers and procedure variables, as illustrated in the following code excerpt of the compiler's symbol table handler (module ORB):

```

1  TYPE Ptrs* = {Pointer, NilTyp};                          (*NilTyp = hidden pointer variable*)
2  Procs* = {Proc, NoTyp};                                   (*NoTyp = hidden procedure variable*)
3
4  PROCEDURE FindHiddenFields(VAR R: Files.Rider; typ: Type; off: LONGINT);
5    VAR fld: Object; i, s: LONGINT;
6  BEGIN
7    IF typ.form IN Ptrs THEN Write(R, Fld); Write(R, 0); Files.WriteNum(R, off)
8    ELSIF typ.form IN Procs THEN Write(R, Fld); Write(R, 0); Files.WriteNum(R, -off-1)
9    ELSIF typ.form = Record THEN fld := typ.dsc;
10   WHILE fld # NIL DO FindHiddenFields(R, fld.type, fld.val + off); fld := fld.next END
11   ELSIF typ.form = Array THEN s := typ.base.size;
12   FOR i := 0 TO typ.len-1 DO FindHiddenFields(R, typ.base, i*s + off) END
13   END
14  END FindHiddenFields;

```

¹⁹ This is acceptable because for record fields, the types *ORB.NilTyp* and *ORB.NoTyp* are not used otherwise.

Assessment of the outlined solution

An obvious shortcoming of the reference checking scheme outlined above is that it requires *additional* run-time information to be present in *all* type descriptors of *all* modules *solely* for the purpose of reference checking, which in turn is *only* needed in the relatively infrequent case of module releases. In addition, modules now also contain an *additional* section in the module block containing the offsets of global procedure variables. However, since there exists only one type descriptor per declared record *type* rather than one per allocated heap *record*, and global procedure variables tend to be rare, the additional memory requirements are negligible²⁰.

Although the operation of *reference checking* may appear expensive at first, in practice there is no performance issue. It is comparable in complexity to garbage collection and thus barely noticeable – at least on systems with small to medium sized dynamic spaces. This is in spite of the fact that for *each* heap object encountered during the *mark* phase *all* modules to be unloaded are checked for references during the *scan* phase. Note, however, that reference checking *stops* when the *first* reference is detected and that module unloading is rare except, perhaps, during development – where however the *number* of references tends to be small.

Thus, the presented solution, which was mainly chosen for its simplicity, appears to be amply sufficient for most practical purposes.

Alternatives considered

Alternative #1:

In the conventional mark-scan scheme outlined above, the initial *mark* phase first marks the dynamic objects reachable by *all other* modules and the *scan* phase then checks for references from the *marked* objects to the modules to be unloaded. In order to improve the efficiency of this scheme, one might be tempted to check for references directly *during* the *mark* phase and simply *stop* marking heap objects as soon as the first reference is found. While this has the potential to be more efficient, it would lead to a number of complications.

First, we note that the pointer rotation scheme used during the *mark* phase temporarily modifies not only the *mark* field, but also the *pointer variable* fields of the encountered heap objects. In Oberon, the Deutsch-Schorr-Waite graph marking algorithm is used to implement this phase²¹. This scheme essentially establishes a *return path* used during future visits of a node, effectively replacing the stack of procedure activations by a stack of marked nodes, i.e. the return path is encoded *in* the heap data structure itself, thereby avoiding recursion and, as a consequence, the consumption of additional stack space during the heap traversal. Thus, one cannot simply exit the *mark* phase when a reference is found, but would also need to undo all pointer modifications made up to that point. The easiest way to achieve this is by letting the *mark* phase run to completion. But this would neutralize the desired performance gain.

²⁰ Adding procedure variable offsets to type descriptors is, strictly speaking, not necessary, as the compiler could always rearrange the memory layout of a record such that all procedure variable fields are placed in a contiguous section at the beginning of it, making their offsets implicit. Thus, it is in principle possible to realize the reference checking scheme without additional fields in type descriptors. We refrained from implementing this refinement for two main reasons. First, it would increase the complexity of the compiler. At first glance, it might seem that in order to rearrange a record's field list, only a small change needs to be made to a single procedure in the compiler (ORP.RecordType in FPGA Oberon). However, a complication arises because records may recursively contain other records, each again containing procedure variables at arbitrary offsets. While it is always possible to "flatten" such recursive record structures, it would make other record operations more complex. For example, assignments to subrecords would become less natural, as their fields would no longer be placed in a contiguous section in memory. Second, the memory savings in type descriptors would be marginal, given that there exists only one type descriptor per record *type* rather than one per allocated heap *record*. Most applications are (or should be) programmed in the conventional programming style, where installed procedures are rare. For example, in the Oberon system, the object-oriented programming style is restricted to the viewer system, which provides distributed control in the form of installed handlers – of which there is usually only one per *type* of viewer. In sum, the benefit obtained by saving a few fields in a relatively small number of type descriptors appears negligible, and therefore the additional effort required to implement this refinement would be hard to justify indeed.

²¹ Herbert Schorr and William M. Waite, "An efficient machine-independent procedure for garbage collection in various list structures", CACM, 10(8):501-506, August 1967.

Second, if one wants to omit in module *Kernel* any reference to the data structure managed by module *Modules*, one would also need to express the *mark* phase as a *generic* heap traversal scheme accepting parametric handler procedures for reference checking, similar to the generic heap *scan* scheme outlined above. Such a generalization would open up the possibility for an erroneous (or malicious) handler procedure to prematurely end the *mark* phase, which in turn may leave the heap in an inconsistent and potentially irreparable state. In sum, one seems well advised not to interfere with the *mark* phase.

Finally, one would still need a *scan* phase to *unmark* the already marked portion of the heap.

Alternative #2:

Another possible variant would be to treat all static and dynamic *procedure variables* and *type tags* like *pointers* (which is what they essentially are) and the *procedures* and *type descriptors* referenced by them like *records* during the *mark* phase of reference checking. This could be achieved by making procedures and type descriptors *look like* records, which in turn could be accomplished by making them carry a *type tag* and a *mark field*.

These additional fields would be inserted as a prefix to (the code section of) *each* declared procedure and (the type descriptor of) *each* declared record type in the module block. All static *procedures* would share a common “procedure descriptor” and all *type descriptors* a common “meta-type descriptor”. These descriptors would contain no tags to extensions and no pointer or procedure variable offsets. Consequently, they can be represented by one and the same shared *global* meta descriptor, which could be stored at a fixed location in the module area for example. The resulting run-time representation of *procedures* and *type descriptors* is shown in Figure 3.

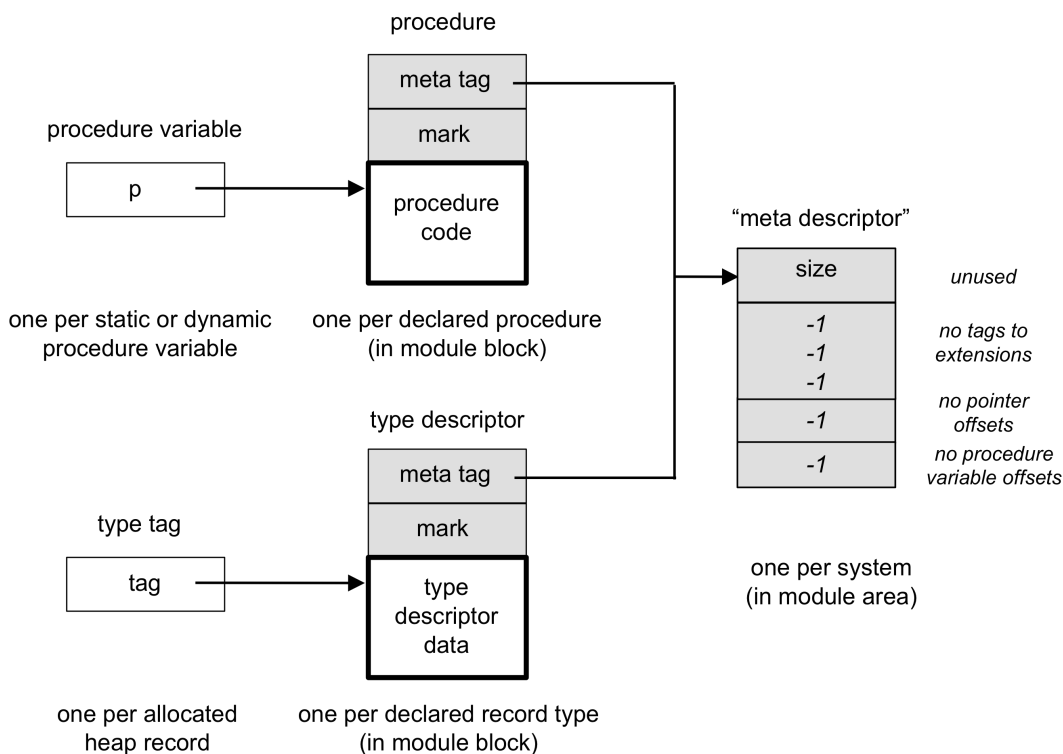


Figure 3: Procedure variable and type tag interpreted as *pointer*, and procedure and type descriptor interpreted as *record*

Alternative #3:

A simpler variant would be to treat procedure variables and type tags as *special cases* during the *mark* phase, eliminating the need for a *meta* tag field as well as the shared *meta* descriptor. The resulting run-time representation of *procedures* and *type descriptors* is shown in Figure 4.

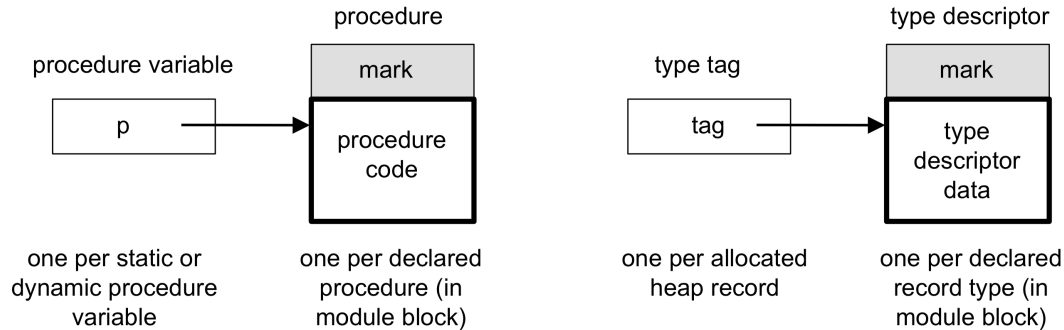


Figure 4: Procedures and type descriptors with an additional *mark* field

With these preparations, the *mark* phase of reference checking could be suitably *extended* to also include *procedure variables* in the list of “pointers” to be traversed. In sum, the mark phase would not only “touch” all reachable *heap* objects, but also those located in the *module blocks*.

While this technique appears appealing at first, a few points are worth pointing out. First, the *extended* mark phase requires an extra *mark* field to be inserted as a prefix to *each* procedure and *each* type descriptor in the module block. Given that most procedures and type descriptors are never referenced, this appears to be overkill. One could therefore decide to add the *mark* field only to *module descriptors* rather than to individual *procedures* or *type descriptors* during the *mark* phase. While this would make the check whether any of the selected modules is referenced a trivial task, it would render the *mark* phase more complex, as it would now need to *locate* the module descriptor pertaining to a given procedure or type descriptor (on systems where additional meta-information is present in the run-time representation of loaded modules, such as the locations of procedures and type descriptors in module blocks or *back pointers* from each object of a module to its module descriptor, the mark phase would be simpler; however, neither FPGA Oberon nor Experimental Oberon offer such *metaprogramming* facilities).

Second, one would still need to mark *all* objects, for the same reason as outlined above, i.e. one cannot simply exit in the middle of the *mark* phase without additional action.

Third, one would also still need an extra *scan* phase to unmark the marked portion of the heap.

Finally, a comparison of the code required to implement the various alternatives showed that our solution is *by far* the simplest: the combined implementation cost of *all* modifications to the runtime representation of type descriptors and descriptors of global module data, the object and symbol file formats and the module loader, is only about 15 source lines of code²², while the *reference checking* phase itself amounts to less than 75 lines (the *total* implementation cost to realize *safe module unloading*, including the ability to unload module *groups* and the automatic collection of no longer referenced hidden modules, amounts to about 250 source lines of code).

²² See *procedures* ORB.InType, ORB.FindHiddenFields, ORG.BuildTD, ORG.Close and Modules.Load.

Appendix: Historical notes on module unloading in the Oberon system

In most Oberon implementations, only *client* references are checked prior to module unloading. *Type* and *procedure variable* references are usually not checked, although various approaches are typically employed to address the case where such references exist. These approaches can be broadly grouped into two main categories: (a) schemes that explicitly allow invalidating references, and (b) schemes where all references must remain unaffected.

a. Schemes that explicitly allow invalidating references

In such schemes, unloading a module from memory *may*, and in general *will*, lead to “dangling” *type* and *procedure variable* references pointing to module data that is no longer valid²³. This includes the important case where a structure rooted in a variable of base type T declared in a base module M contains elements of an extension T' defined in a client module M'. Such elements typically contain both *type* references (type tags) and *procedure variable* references (handler procedures) referring to module M'. This is common in the Oberon viewer system, for example, where M is module Viewers and M' may be a graphics editor. If the client module M' is unloaded and the module block previously occupied by it is overwritten by another module loaded later, any still existing references to M' become *invalid* at that moment. Global procedure variables declared in other modules may also refer to procedures declared in module M', although this case is much less common (global procedure variables tend to be used mainly for procedures declared in the same module).

A variety of approaches have been employed in different implementations of the Oberon system to cope with the introduced dangling *type* or *procedure variable* references:

1. On systems that use a *memory management unit* to perform virtual memory management, such as Ceres-1 or Ceres-2, one approach is to *unmap* the module space of an unloaded module from virtual memory, thereby *invalidating* future references to it. When this is done, any still existing *type* or a *procedure variable* references point to a now *unallocated* page, and consequently any attempt to access such a page later will result in a system *trap*.

We consider this an unfortunate proposal for several reasons. First, users generally have no way of knowing *whether* it is in fact safe to unload a module, yet they are allowed to do so. Second, after having unloaded it, they still don't know whether references from other loaded modules still exist – until a *trap* occurs. But then it is usually too late. While the trap itself will actually *prevent* a system crash (as intended), the user may *still* need to restart the system in order to recover an environment without any “frozen” parts, for example displayed viewers that may have been opened by the unloaded module. Note also that this solution requires special hardware support, which may not be available on all systems.

2. On systems that do *not* use virtual memory, such as Ceres-3 or FPGA Oberon on RISC, the easiest way to cope with dangling references is to simply *ignore* them, i.e. to *always* release the memory associated with a module without any further precautions (unless clients exist). While an attempt to access an unloaded module goes undetected *initially*, the system is still left in an *unsafe* state, which will become *unstable* the moment another module loaded later *overwrites* the previously released module block *and* other loaded modules still refer to its *type descriptors* or *procedures*. This is of course undesirable.

²³ If the programming language Oberon-2 is used, there can also be references to method tables (which however are typically allocated within type descriptors).

3. Another approach, which however can be used only for *procedure variable* references, is to identify *all* such references and make them refer to a *dummy* procedure, thereby preventing a run-time error when such “fixed up” procedures are called later. This solution requires the system to *know* the locations of all procedure variables in all static and dynamic objects at run time. It was used in an earlier version of Experimental Oberon, but was later discarded, mainly because the resulting effect on the *overall* behavior of the system is essentially impossible to predict (or even detect) by the user. The fact that *some* procedure variables *somewhere* in the system no longer refer to *real* but to *dummy* procedures typically becomes “visible” only through the *absence* of some action – such as mouse tracking if the unloaded module contained a viewer handler, for instance.
4. On systems that use *indirection* for procedure calls via a so-called “link table”, the same effect can be achieved by setting the *link table entries* for all referenced *procedures* of the module to be unloaded to *dummy* entries, rather than locating and modifying each individual procedure *call* that may exist anywhere in the system.

Note that using a link table to implement indirection for procedure calls is viable only on systems with *efficient* hardware support for it. On such systems, an “address” of a procedure is not a real memory address, but an *index* to this translation table – which the caller consults for every procedure *call*, in order to obtain the actual memory location of the called procedure. Indirection for procedure calls via a link table was used in some of the earlier versions of the Oberon system on Ceres computers, which were based on the (now defunct) NS32000 processor. This processor featured a *call external procedure* instruction (CXP *k*, where *k* is the index of the link table entry of the called procedure), which sped up the process of calling external procedures significantly²⁴. Later versions of the NS processor, however, internally re-implemented the *same* instruction using microcode, which negatively impacted its performance²⁵. For this and a variety of other reasons, the CXP *k* instruction – and with it the *link table* – were no longer used in later versions of the Oberon system.

5. For *type* references, it is actually possible to determine at *compile* time, whether a module *may* lead to references from other loaded modules at *run* time. The criteria is the following: if a module *M'* does *not* declare record types which are extensions *T'* of an imported type *T*, then records declared in *M'* *cannot* ever be inserted in a data structure rooted in a variable *v* of the imported type *T* – precisely *because* they are not extensions of *T* (in the Oberon programming language, an assignment *p* := *p'* is allowed only if the type of *p'* is the same as that of *p* or an extension of it). One *could* therefore introduce a rule that a module *M'* can be safely dispensed *only* if it does *not* declare record types, which are extensions *T'* of an imported type *T*. The flip side of such a rule, however, is that modules that actually declare such types can *never* be unloaded – unless, of course, other ways to safely unload such modules are implemented. Also, procedure variable references are *not* covered by this rule.

Even though most of these approaches have actually been realized in various implementations of the Oberon system, we consider none of them truly satisfactory. In our view, these schemes appear to only tinker with the symptoms of a problem that would not exist, if only one adopted the rule to *disallow* the removal (from memory) of still referenced module data.

²⁴ The use of the link table also increased code density considerably (as only 8 bits for the index instead of 32 bits for the full address were needed to address a procedure in every procedure call). In addition, the link table used by the CXP instruction allowed for an expedient linking process at load time (as there are far fewer conversions to be performed by the module loader – one for every referenced procedure instead of one for every procedure call) and also eliminated the need for a fixup list (list of the locations of all external procedure references to be fixed up by the module loader) in the object file. A disadvantage is, of course, the need for a (short) link table.

²⁵ The internal re-implementation of the CXP instruction using microcode in later versions of the NS processor followed the general industry trend of implementing only frequent, simple instructions directly with hardware, while interpreting more complex instructions using internal microcode. In general, with the advent of highly regular reduced instruction set computers (RISCs) in the mid 1980s and early 1990s, the trend towards offering microprocessors providing a smaller set of simple instructions, most of them executing in a single clock cycle, combined with fairly large banks of (fast) registers, continued (and does so to this date).

The main issue is that the moment one *allows* modules to release their associated memory in spite of references to them from other loaded modules, the resulting *dangling* references must be dealt with *somehow* in order to prevent an almost certain system *crash*. However, *fixing up* or *invalidating* references will *always* remove essential information from the system. As a consequence, the run-time behavior of the modified system becomes *essentially unpredictable*, as other loaded modules may *critically* depend on the removed functionality. For example, unloading a module that contains an installed handler procedure of a *contents frame* may render it impossible to *close* the enclosing *menu viewer* that contains it, thereby leading to a system with “frozen” parts. A similar problem may occur with references to *type descriptors*, if they are not persisted in memory *after* unloading their associated modules.

b. Schemes where all references must remain unaffected

The second possible interpretation of *unloading* a module consists of schemes where all references *must* remain unaffected at all times. In such schemes, module unloading can be viewed as an implicit mandate to preserve “critical module data”, as long as references to the unloaded module still exist. Various possible ways to fulfill this mandate exist:

1. One could of course simply exit the *unload* command with an error message, whenever such references are detected. The user, however, may then be “stuck” with modules that he can *never* unload because they are referenced by modules over which he has no explicit control.
2. But the mandate could also be fulfilled by persisting any still referenced module data to a “safe” location before unloading the associated module.

For *type descriptors* referenced by *type tags* in dynamically allocated heap records a simple solution exists: allocate them *outside* the module blocks in order to persist them beyond the lifetime of their associated module. One possibility is to allocate them in the *heap* itself at module load time²⁶. This method has been implemented in Ceres-Oberon, for example. It eliminates dangling *type* references altogether and therefore also the *need* to check for them at run time, since type descriptors are now unaffected by module unloading.

For *procedures* referenced by *procedure variables* in either static or dynamically allocated objects no such simple solution exists. Unless one accepts to *invalidate* or *fix up* procedure variable references (as outlined in the previous section), the only way to “persist” procedures is to persist the *entire* module (recall that procedures may *access* global module data or *call* other procedures declared in the same module).

We conclude that if one wants to handle *type* and *procedure variable* references, one *cannot* unload a module block from memory, as long as references to it still exist²⁷.

3. One way to automatically persist type descriptors *and* procedures consists of simply *never* releasing the module block of a module once it is loaded. Instead, when the user requests the unloading of a module, it is only removed from the *list* of loaded modules. This amounts to *renaming* the module, with the implication that a newer version of the same module (with the same name) can be reloaded again²⁸. This approach has been chosen in MacOberon²⁹.

²⁶ Note that one cannot simply move type descriptors around in memory, as their addresses are (typically) used to implement type tests and type guards. By allocating them in the heap at module load time, one avoids the need to move them to a different location when a module is unloaded.

²⁷ Of course, “mixed” variants are also possible. For example, one could allocate type descriptors in the heap at module load time (as in Ceres-Oberon), and either fix up all procedure variable references or prevent the release of a module block if such references still exist; however, most modules referenced by type tags are *also* referenced by procedure variables – this is in fact the typical case for dynamic records containing installed handler procedures. Thus, it seems more natural to employ the *same* approach for both type and procedure variable references.

²⁸ In a specific implementation, one might choose to make the module completely anonymous or modify the module name such that one can no longer import it (e.g., by inserting an asterisk).

²⁹ <http://e-collection.library.ethz.ch/eserv/eth:3269/eth-3269-01.pdf> (The Implementation of MacOberon, 1990)

Since the associated memory of a module is never released, the issue of dangling type or procedure variable references is avoided altogether, as they simply cannot exist. However, it can also lead to higher-than-necessary memory usage if a module is repeatedly loaded and unloaded (typical during *development*). Nevertheless, such an approach may be viewed as adequate on *production* systems, where module unloading is rare.

4. A *refinement* of the approach outlined above consists of *initially* removing a module only from the module list (as in MacOberon), but *in addition* releasing its associated memory *as soon as* there are no more type or procedure variable references to it. If this is done in an automatic fashion (for example as part of a background process), module data is truly “kept in memory for exactly as long as necessary and removed from it as soon as possible”.

This is the approach chosen in Experimental Oberon. A variation of it was used in some versions of SparcOberon³⁰.

In sum, module unloading schemes where references remain *unaffected* at all times avoid many of the complications that are inherent in schemes that explicitly *allow* invalidating references. A (small) price to pay is to keep modules loaded in memory, as long as references to them exist.

However, on *production* systems, there is typically no need to keep multiple copies of the same module in memory, while on *development* systems it is totally acceptable.

Finally, we note that on modern computers the amount of available memory, and therefore also the amount of *dynamic* data that may be allocated by modules, typically far exceeds the size of the module blocks holding the program code and global variables. Hence, not releasing module blocks immediately after a module *unload* operation typically has a rather negligible impact on overall memory usage.

* * *

³⁰ <http://e-collection.library.ethz.ch/eserv/eth:7103/eth-7103-01.pdf> (SPARC-Oberon User's Guide and Implementation, 1990/1991)