

Safe module unloading

in the Oberon operating system

Andreas Pirklbauer

1.7.2019

Purpose

This technical note describes a possible implementation of *safe module unloading* for the Oberon operating system, as realized in Experimental Oberon¹, a revision of FPGA Oberon².

Module unloading in the Oberon operating system

In the Oberon system, there exist three possible types of references to a loaded module M³:

1. *Client references* exist when other loaded modules *import* module M.
2. *Type references* exist when type tags (addresses of type descriptors) in *dynamic* objects reachable by other loaded modules refer to descriptors of types *declared* in module M.
3. *Procedure variable references* exist when procedure variables in *static* or *dynamic* objects reachable by other loaded modules refer to procedures *declared* in module M.

In most Oberon implementations, only *client* references are checked prior to module unloading, i.e. if clients exist among the other loaded modules, a module or module group is not unloaded from the system. *Type* and *procedure variable* references are usually not checked, although various approaches can be employed to address the case where such references exist⁴.

As a result, module unloading has traditionally been *unsafe* in the Oberon operating system, as removing a module from memory *may*, and in general *will*, lead to “dangling” references that point to module data that is no longer valid. This means that the system will become *unstable* the very moment another module loaded later *overwrites* a previously released module block *and* other loaded modules still refer to its *type descriptors* or *procedures*.

¹ <http://www.github.com/andreaspirklbauer/Oberon-experimental>

² <http://www.projectoberon.com>

³ An Oberon module can be viewed as a container of types, variables and procedures. Types can be declared *global* (in which case they can be exported and referenced by name in client modules) or *local* to a procedure (in which case they cannot be exported). Variables can be declared as *global* variables (allocated in the module area when a module is loaded) or as *local* variables (allocated on the stack when a procedure is called). Anonymous variables with no explicit name declared in the program can be dynamically allocated in the heap via the predefined procedure *NEW*. Procedures can be declared as *global* or *local* procedures, and can be assigned to procedure variables. Thus, in general there can be type, variable, procedure and procedure variable references from static or dynamic objects of other modules to static or dynamic objects of the modules to be unloaded. However, only *dynamic* type references and *static* and *dynamic* procedure variable references need to be checked during module unloading for the following reasons: First, *static* type, variable or procedure references from other modules can only refer by *name* to types, variables or procedures *declared* in the modules to be unloaded. Such references are already handled via their import/export relationship during module unloading (if clients exist, a module or module group is never unloaded) and therefore don't need to be checked separately. Second, *dynamic* pointer references from global or dynamic *pointer* variables of other modules to *dynamic* objects reachable by the modules to be unloaded *should* not be checked, as they should not prevent module unloading. In the Oberon system, such references will be handled by the garbage collector during a future garbage collection cycle, i.e. heap records reachable by the just unloaded modules *and* other still loaded modules will not be collected, whereas heap records that *were* reachable *only* by the unloaded modules *will* be collected – as they should. Thus, the handling of pointer references is delegated to the garbage collector. Finally, *pointer* variable references to *statically* declared objects are only possible by resorting to low-level facilities and should be avoided – and, in fact, be disallowed (pointers should point exclusively to *anonymous* variables allocated when needed during program execution).

⁴ See the appendix for historical notes on module unloading in the Oberon operating system.

Implementing safe module unloading

In Experimental Oberon, a revision of FPGA Oberon⁵, *all* possible types of references to a loaded module or module group are checked as follows prior to module *unloading* (Figure 1):

- If clients exist among the loaded modules, a module or module group is never unloaded.
- If no client, type or procedure variable references to a module or module group exist in the remaining modules or data structures, it is unloaded and its associated memory is released.
- If no clients, but type or procedure variable references exist, the module *unload* command takes no action (by default) and merely displays the names of the modules containing the references that caused the removal to fail. If, however, the *force* option */f* is specified⁶, the modules are initially removed (only) from the *list* of loaded modules, without releasing their associated memory. Such “hidden” modules are later physically removed from *memory*, as soon as there are no more references to them from anywhere in the system.

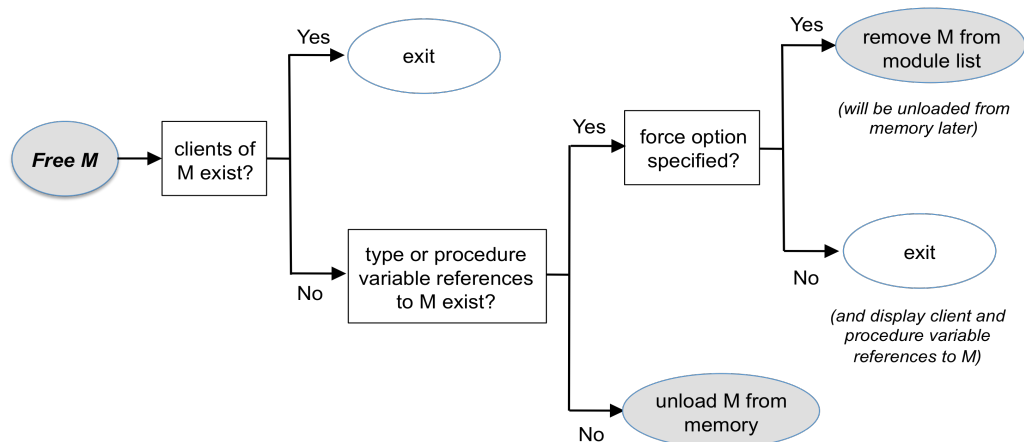


Figure 1: Safe module unloading in Experimental Oberon

Removing a module only from the *list* of loaded modules amounts to *renaming* it, with the implication that a newer version of the same module with the same name can be reloaded again, without having to unload (from memory) earlier versions that are still referenced⁷. Unloading a module from *memory* frees up the memory area previously occupied by the module block⁸. To automate the unloading of no longer referenced *hidden* module data, a command *Modules.Collect* has been included in the background task handling garbage collection. It checks all possible combinations of *k* modules chosen from *n* hidden modules for references to them, and removes those module subgroups from memory that are no longer referenced. The tool command *System.Collect* also invokes *Modules.Collect*.

In sum, module unloading does not affect any references, as module data is kept in memory for exactly as long as necessary and removed from it as soon as possible.

⁵ <http://www.projectoberon.com>

⁶ The force option */f* must be specified at the *e n d* of the list of modules to be unloaded, e.g., *System.Free M1 M2 M3/f*

⁷ Modules removed only from the list of loaded modules, but not from memory, are marked with an asterisk in the output of the command *System.ShowModules*. Commands of such “hidden” modules can be accessed by either specifying their module number or their (modified) module name, both of which are displayed by the command *System.ShowModules*. In both cases, the corresponding command text must be enclosed in double quotes. If a module *M* carries module number 14, for instance, one can activate a command *M.P* also by clicking on the text “14.P”. Typical use cases include hidden modules that still have background tasks installed which can only be removed by activating a command of the hidden modules themselves, or hidden modules that still have open viewers. If the command to close a viewer is displayed in the viewer’s menu bar, the user can manually edit the command text (by clicking within its bottom two pixel lines), replace the module name by its module number and enclose the modified command text in double quotes. Although this is somewhat clumsy, it at least enables the user to close the viewer. An alternative approach is to provide a “Close” command that also accepts the marked viewer as argument (using procedure *Oberon.MarkedViewer*).

⁸ In FPGA Oberon 2013 on RISC and in Experimental Oberon, the module block includes the module’s type descriptors. In some other Oberon implementations, such as Oberon on Ceres, type descriptors are not stored in the module block, but are dynamically allocated in the heap at module load time, in order to persist them beyond the lifetime of their associated modules. In Experimental Oberon, no such special precaution is necessary, as module blocks are removed only from the list of loaded modules and not from memory, if they are still referenced by other modules. Thus, type descriptors can safely be stored in the (static) module blocks.

For example, older versions of a module's code can still be executed if they are referenced by static or dynamic procedure variables in other modules, even if a newer version of the module has been reloaded in the meantime⁹. Type descriptors also remain accessible to other modules for exactly as long as needed. This covers the important case where a structure rooted in a variable of base type *T* declared in a base module *M* contains elements of an extension *T'* defined in a client module *M'*, which is unloaded¹⁰. Such elements typically contain both *type* references (type tags) and *procedure variable* references (handler procedures) referring to *M'*. This is common in the Oberon viewer system, for example, where *M* is module *Viewers*.

If a module *group* is to be unloaded and there exist references *only* within this group, it is unloaded *as a whole*. This can be used to remove module groups with *cyclic* references¹¹. It is also possible to release *entire subsystems* of modules. The command *System.FreeImports* attempts to unload the specified modules and all their direct and indirect *imports*. It may be used for conventional module stacks with a single top module, for example a compiler. By contrast, the command *System.FreeClients* unloads the specified modules and all their direct and indirect *clients*. This variant may be used for module stacks written in the object-oriented programming style with their typically inverted module hierarchy, for example a graphics editor.

Note that these unloading strategies amount to *heuristics* and may tend to unload rather large parts of the system, which may not be desired. Thus, the recommended way to unload modules is to use the command *System.Free* with a *specific* set of modules provided as parameters. For added convenience, the tool commands *System.ShowRefs* and *System.ShowGroupRefs* can be used to identify all modules containing references to a given module or module group.

Implementation aspects

To unload a module group, the module *unload* command *selects* the modules to be unloaded using procedure *Modules.Select* and then invokes procedure *Modules.Check*. Clients are checked first. This is done by verifying whether *unselected* modules import *selected* modules¹²:

```

1  PROCEDURE FindClients*(proc: ImpHandler; VAR res: INTEGER);
2    VAR mod, imp, m: Module; p, q: INTEGER; continue: BOOLEAN;
3  BEGIN res := noref; m := root; continue := proc # NIL;
4    WHILE continue & (m # NIL) DO
5      IF (m.name[0] # 0X) & m.selected & (m.refcnt > 0) THEN mod := root;
6        WHILE continue & (mod # NIL) DO
7          IF (mod.name[0] # 0X) & ~mod.selected THEN p := mod.imp; q := mod.cmd;
8            WHILE p < q DO imp := Mem[p];
9              IF imp = m THEN INC(res, proc(mod, imp, continue)); p := q ELSE INC(p, 4) END
10           END
11         END ;
12         mod := mod.next
13       END
14     END ;
15     m := m.next
16   END
17 END FindClients;
```

⁹ If an older version of a module's code accesses global variables (of itself or of other modules), it will automatically access the "right" version of such variables – as it should.

¹⁰ For type references, it is possible to determine at compile time, whether a module may lead to references from other modules at run time. The criteria is the following: if a module *M'* does not declare record types which are extensions *T'* of an imported type *T*, then records declared in *M'* cannot be inserted in a data structure rooted in a variable *v* of the imported type *T* – precisely because they're not extensions of *T* (in the Oberon programming language, an assignment *p := p'* is allowed only if the type of *p'* is the same as that of *p* or an extension of it).

¹¹ In Oberon, cyclic references can be created by pointers or procedure variables, whereas cyclic module imports are normally not allowed. However, through a tricky sequence of compilation and editing steps, it is in fact possible to *construct* cyclic module imports, which cannot be detected by the compiler. But even though they can be created, such modules are not allowed to be loaded, as the module loader of Experimental Oberon – adopting the approach chosen in Original Oberon and FPGA Oberon – would enter an endless recursion, eventually leading to a stack overflow or out-of-memory condition – a totally acceptable solution for such an artificially constructed case. But even if modules with cyclic module imports *were* allowed to be loaded, Experimental Oberon would handle them correctly upon unloading, i.e. if no external clients or references exist, such a module group would simply be unloaded as a whole – as it should.

¹² *Mem* stands for the entire memory and assignments involving *Mem* are expressed as *SYSTEM.GET(a, x)* for *x := Mem[a]* and *SYSTEM.PUT(a, x)* for *Mem[a] := x*.

If clients exist among the loaded modules, no further action is taken and the *unload* command exits. If no clients exist, *type* and *procedure variable* references are checked next. References from *dynamic* objects in the heap are checked using a conventional *mark-scan* scheme:

```

1  PROCEDURE FindDynamicRefs*(type, proc: RefHandler; VAR resType, resProc: INTEGER; all: BOOLEAN);
2    VAR mod: Module;
3    BEGIN mod := root;
4      WHILE mod # NIL DO
5        IF (mod.name[0] # 0X) & ~mod.selected THEN Kernel.Mark(mod.ptr);
6          IF ~all THEN Kernel.Scan(type, proc, mod.name, resType, resProc) END
7        END ;
8        mod := mod.next
9      END ;
10     IF all THEN Kernel.Scan(type, proc, "", resType, resProc) END
11  END FindDynamicRefs;

```

During the initial *mark* phase, heap records reachable by all *named* global pointer variables of *all other* loaded modules are marked (line 5), thereby excluding records reachable *only* by the specified modules themselves. This automatically recognizes module *groups* and ensures that when a module or module group is referenced *only* by itself, it can still be unloaded. The *scan phase* (line 6 or 10), implemented as a separate procedure *Scan*¹³ in module *Kernel*, scans the heap sequentially, unmarks all *marked* records and checks whether the *type tags* of the marked records point to descriptors of types declared in the *selected* modules, and whether *procedure variables* declared in these heap records refer to procedures declared in those same modules.

An additional boolean parameter *all* allows the caller to specify whether the *mark* phase should first mark all heap records that are reachable by *all* other loaded modules before initiating the *scan* phase *once* (used for module unloading), or whether the *scan* phase should be initiated for each module separately, after marking the heap records reachable by it (used for enumerating the references to each module *individually*).

Finally, the check for *procedure variable* references is also performed for all *global* procedure variables, whose offsets in the module's data section are obtained from an array in the module's *meta* data section, headed by the link *mod.pvr* in the module descriptor (see below):

```

1  PROCEDURE FindStaticRefs*(proc: RefHandler; VAR res: INTEGER);
2    VAR mod: Module; pref, pvadr, r: LONGINT; continue: BOOLEAN;
3    BEGIN res := noref; mod := root; continue := proc # NIL;
4      WHILE continue & (mod # NIL) DO
5        IF (mod.name[0] # 0X) & ~mod.selected THEN
6          pref := mod.pvr; pvadr := Mem[pref];
7          WHILE continue & (pvadr # 0) DO r := Mem[pvadr];
8            INC(res, proc(pvadr, r, mod.name, continue));
9            INC(pref, 4); pvadr := Mem[pref]
10        END
11      END ;
12      mod := mod.next
13    END
14  END FindStaticRefs;

```

Note that the procedures *FindClients*, *FindDynamicRefs* and *FindStaticRefs* are expressed as *generic* object traversal schemes which accept *parametric* handler procedures that are called for each encountered object. This allows these procedures to be used for *other* purposes as

¹³ The original procedure *Kernel.Scan* (implementing the scan phase of the Oberon garbage collector) has been renamed to *Kernel.Collect*, in analogy to procedure *Modules.Collect*.

well, for example to *count* or *enumerate* the client, type or procedure variable references to *any* given (loaded) module or module group.

In order to omit in module *Kernel* any reference to the module list rooted in module *Modules*, procedure *Kernel.Scan* is also expressed as a *generic* heap scan scheme. The following is a simplified version of this scheme¹⁴:

```

1  PROCEDURE Scan*(type, proc: Handler; s: ARRAY OF CHAR; VAR resType, resProc: INTEGER);
2  VAR offadr, offset, p, r, mark, tag, size: LONGINT; continue: BOOLEAN;
3  BEGIN p := heapOrg; resType := 0; resProc := 0; continue := (type # NIL) OR (proc # NIL);
4  REPEAT mark := Mem[p+4];
5  IF mark < 0 THEN (*free*) size := Mem[p]
6  ELSE (*allocated*) tag := Mem[p]; size := Mem[tag];
7  IF mark > 0 THEN (*marked*) Mem[p+4] := 0;
8  IF continue THEN
9  IF type # NIL THEN INC(resType, type(p, tag, s, continue)) END ;
10 IF continue & (proc # NIL) THEN offadr := tag + 16; offset := Mem[offadr];
11 WHILE offset # -1 DO (*skip pointers*) INC(offadr, 4); offset := Mem[offadr] END ;
12 INC(offadr, 4); offset := Mem[offadr];
13 WHILE continue & (offset # -1) DO (*procedures*) r := Mem[p+8+offset];
14 INC(resProc, proc(p+8+offset, r, s, continue));
15 INC(offadr, 4); offset := Mem[offadr]
16 END
17 END
18 END
19 END
20 END ;
21 INC(p, size)
22 UNTIL p >= heapLim
23 END Scan;
```

This scheme calls *parametric* handler procedures for individual elements of each *marked* heap record rather than *directly* checking whether these records contain type or procedure variable references to the modules to be unloaded. Procedure *type* is called with the memory address of the *heap record* itself as the *source* of the reference followed by the *type tag* (address of the type descriptor) as its *destination* (line 9). Procedure *proc* is called for each procedure variable declared in the same record with the address of the *procedure variable* as the source followed by the address of the referenced *procedure* as the destination (line 14). The results of these handler calls are *separately* added up for each handler procedure and returned in the variable parameters *resType* and *resProc*.

An additional variable parameter *continue* allows the handler procedures to indicate to the caller that they are no longer to be called (lines 8, 10, 13). The scan process itself continues, but only to *unmark* the remaining marked records (line 7).

Procedure *Modules.Check* uses procedures *FindClients*, *FindDynamicRefs* and *FindStaticRefs* by passing its private handler procedures *HandleClient* and *HandleRef*.

We emphasize that procedure *Modules.Check* is not only called when a module is unloaded by the *user*, but also by the Oberon *background task* that removes no longer referenced *hidden* module data from memory. Thus, it must be written such that it can correctly handle *both* visible *and* hidden modules in the module data structure rooted in module *Modules*.

¹⁴ The simplified version scans only *record* blocks allocated via *NEW(p)*, where *p* is a *POINTER TO RECORD*. The full implementation also covers *array* blocks allocated in the heap.

Implementation prerequisites

In order to make the outlined validation pass possible, type descriptors of *dynamic* objects¹⁵ allocated in the *heap* as well as the descriptors of *global* module data located in *static* module blocks contain a list of *procedure variable offsets*, adopting an approach employed in one of the earlier implementations of the Original Oberon system (MacOberon)¹⁶. The resulting run-time representation of a *dynamic* record and its associated type descriptor is shown in Figure 2 (the *method table* used to implement Oberon-2 style *type-bound procedures* is not discussed here).

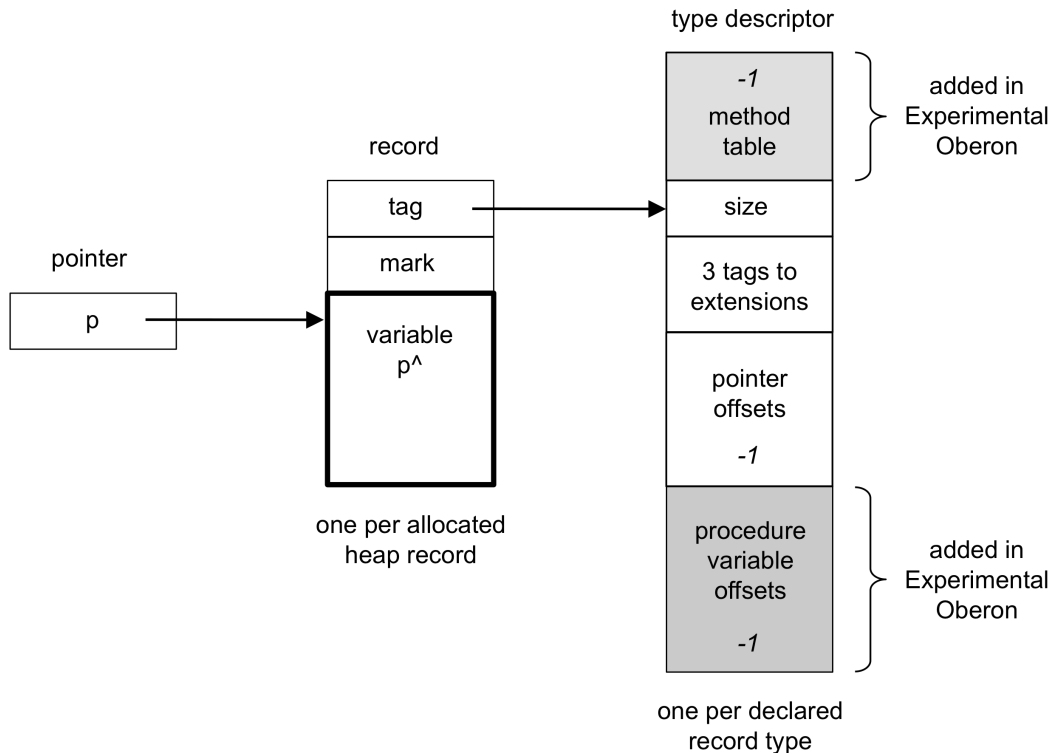


Figure 2: Run-time representation of a dynamic record and its type descriptor in Experimental Oberon

The descriptors also contain the offsets of *hidden* (not exported) procedure variables, enabling the module *unload* command to check *all* possible procedure variable locations in the entire system for possible procedure variable references to the modules to be unloaded.

To make the offsets of hidden *procedure variables* in exported record types available to client modules, symbol files also include them. An importing module may, for example, declare a global variable of an imported record type, which contains *hidden* procedure variables, or declare a record type that contains or extends an imported one. We recall that in Original Oberon, hidden *pointers*, although not exported and therefore invisible in client modules, are already included in symbol files because their offsets are needed for garbage collection¹⁷. Similarly, in Experimental Oberon, the locations of visible *and* hidden procedure variables are needed for reference checking during module unloading, as shown in the following example.

¹⁵ See chapter 8.2, page 109, of the book *Project Oberon 2013 Edition* for a detailed description of an Oberon type descriptor. A type descriptor contains certain information about dynamically allocated records that is shared by all allocated objects of the same record type (such as its size, information about type extensions and the offsets of pointer fields within the record).

¹⁶ <http://e-collection.library.ethz.ch/eserv/eth:3269/eth-3269-01.pdf> (The Implementation of MacOberon, 1990)

¹⁷ See chapter 12.6.2, page 41, of the book *Project Oberon 2013 Edition*.

```

1  MODULE M0;
2  TYPE Proc* = PROCEDURE; Rec* = RECORD p*, q: Proc END ;  (*q is a hidden field*)
3  PROCEDURE Q*(r: Rec); BEGIN r.q END Q;
4  PROCEDURE Init*(VAR r: Rec; p, q: Proc); BEGIN r.p := p; r.q := q END Init;
5  END M0.
6
7  MODULE M1;
8  IMPORT M0;
9  VAR r: M0.Rec;
10 PROCEDURE Init*(p, q: M0.Proc); BEGIN M0.Init(r, p, q) END Init;
11 END M1.
12
13 MODULE M2;
14 IMPORT M1;
15 PROCEDURE Q; BEGIN END Q;
16 PROCEDURE Set1*; BEGIN M1.Init(NIL, NIL) END Set1;
17 PROCEDURE Set2*; BEGIN M1.Init(NIL, Q) END Set2;
18 END M2.
19
20 M2.Set1 System.Free M2 ~      ... can unload M2
21 M2.Set2 System.Free M2 ~      ... can't unload M2 (as M2.Q is referenced in global procedure variable M1.r.q)

```

Here the global record r declared in module $M1$ contains a hidden field $M1.r.q$ that is accessible only in the exporting module $M0$. If another module $M2$ installs a procedure $M2.Q$ in this hidden record field, a procedure variable reference from $M1$ to $M2$ is created, with the (intended) effect that module $M2$ can no longer be unloaded. In order to be able to check for such references at run time, the record field offset of the field $M1.r.q$ must be available when compiling $M1$ ¹⁸.

Hidden *pointers* are included in symbol files without their names, and their (imported) type in the symbol table is of the form $ORB.NilTyp$. Similarly, hidden *procedure variables* are also included in symbol files without their names, but their (imported) type in the symbol table is of the form $ORB.NoTyp$ ¹⁹. Additional details are to be looked up in the procedures $ORB.FindHiddenFields$, $ORG.FindRefFlds$, $ORG.NofRefs$ and $ORG.FindRefs$.

Assessment of the outlined solution

An obvious shortcoming of the reference checking scheme outlined above is that it requires *additional* run-time information to be present in *all* type descriptors of *all* modules *solely* for the purpose of reference checking, which in turn is *only* needed in the relatively infrequent case of module releases. In addition, modules now also contain an *additional* section in the module block containing the offsets of global procedure variables. However, since there exists only one type descriptor per declared record *type* rather than one per allocated heap *record*, and global procedure variables tend to be rare, the additional memory requirements are negligible²⁰.

Although the operation of *reference checking* may appear expensive at first, in practice there is no performance issue. It is comparable in complexity to garbage collection and thus barely noticeable – at least on systems with small to medium sized dynamic spaces (heaps). This is in

¹⁸ The compiler will include the record field offset in the object file of $M1$, from where the module loader will transfer it to the corresponding array in the module's meta data section of $M1$.

¹⁹ This is acceptable because for record fields, the types $ORB.NilTyp$ and $ORB.NoTyp$ are not used otherwise.

²⁰ Adding procedure variable offsets to type descriptors is, strictly speaking, not necessary, as the compiler could always rearrange the memory layout of a record such that all procedure variable fields are placed in a contiguous section at the beginning of it, making their offsets implicit. Thus, it is in principle possible to realize the reference checking scheme without additional fields in type descriptors. We refrain from implementing this refinement for two main reasons. First, it would increase the complexity of the compiler. At first glance, it might seem that in order to rearrange a record's field list, only a small change needs to be made to a single procedure in the compiler ($ORP.RecordType$ in *FPGA Oberon 2013*). However, a complication arises because records may recursively contain other records, each again containing procedure variables at arbitrary offsets. While it is always possible to "flatten out" such recursive record structures, it would make other record operations more complex. For example, assignments to subrecords would become less natural, as their fields would no longer be placed in a contiguous section in memory. Second, the memory savings in type descriptors would be marginal, given that there exists only one type descriptor per declared record *type* rather than one per allocated heap *record*. Most applications are programmed in the conventional programming style, where installed procedures are rare. For example, in the Oberon operating system, the object-oriented programming style is restricted to the viewer system, which provides distributed control in the form of installed handlers – of which there is usually only one per *type* of viewer. In sum, the benefit obtained by saving a few fields in a relatively small number of type descriptors appears negligible, and therefore the additional effort required to implement this refinement would be hard to justify.

spite of the fact that for *each* heap record encountered in the *mark* phase *all* modules to be unloaded are checked for references during the *scan* phase. Note, however, that reference checking *stops* when the *first* reference is detected, and that module unloading is rare except, perhaps, during development – where however the *number* of references tends to be small.

Thus, the presented solution, which was mainly chosen for its simplicity, appears to be amply sufficient for most practical purposes.

Alternatives considered

Alternative #1:

In the conventional mark-scan scheme outlined above, the initial *mark* phase first marks the heap objects reachable by *all other* loaded modules and the *scan* phase subsequently checks for references from the *marked* records to the modules to be unloaded. In order to further improve the efficiency of this scheme, one might be tempted to check for references directly *during* the *mark* phase and simply *stop* marking records as soon as the first reference is found. While this has the potential to be more efficient, it would lead to a number of complications.

First, we note that the pointer rotation scheme used during the *mark* phase temporarily modifies not only the *mark* field, but also the *pointer variable* fields of the encountered heap records. In the Oberon operating system, the Deutsch-Schorr-Waite graph marking algorithm²¹ is used for this purpose. This scheme essentially establishes a *return path* used during future visits of a node, effectively replacing the stack of procedure activations by a stack of marked nodes. Thus, one cannot simply exit the *mark* phase when a reference is found, but would also need to undo all pointer modifications made up to that point. The easiest way to achieve this is by letting the *mark* phase run to completion. But this would neutralize the desired performance gain.

Second, if one wants to omit in module *Kernel* any reference to the data structure managed by module *Modules*, one would also need to express the *mark* phase as a *generic* heap traversal scheme accepting parametric handler procedures for reference checking, similar to the generic heap *scan* scheme outlined above. While this would be straightforward to implement, such a generalization would open up the possibility for an erroneous handler procedure to prematurely end the *mark* phase, which in turn may leave the heap in an inconsistent and potentially irreparable state. In sum, one seems well advised not to meddle with the *mark* phase.

Finally, one would still need a *scan* phase to *unmark* the already marked portion of the heap.

Alternative #2:

Another possible variant would be to treat *procedure variables* and *type tags* like *pointers*, and *procedures* and *type descriptors* referenced by them like *records* during the *mark* phase of reference checking. This could be achieved by making procedures and type descriptors *look like* records, which could be accomplished by making them carry a *type tag* and a *mark field*.

These additional fields would be inserted as a prefix to (the code section of) *each* declared procedure and (the type descriptor of) *each* declared record type in the module block. All static *procedures* would share a common “procedure descriptor” and all *type descriptors* a common “meta-type descriptor”. These descriptors would contain no tags to extensions and no pointer

²¹ Herbert Schorr and William M. Waite, “An efficient machine-independent procedure for garbage collection in various list structures”, CACM, 10(8):501-506, August 1967.

offsets. Consequently, they can be represented by one and the same shared global “meta descriptor”, which could be stored at a fixed location within the module area for example. The resulting run-time representation of *procedures* and *type descriptors* is shown in Figure 3.

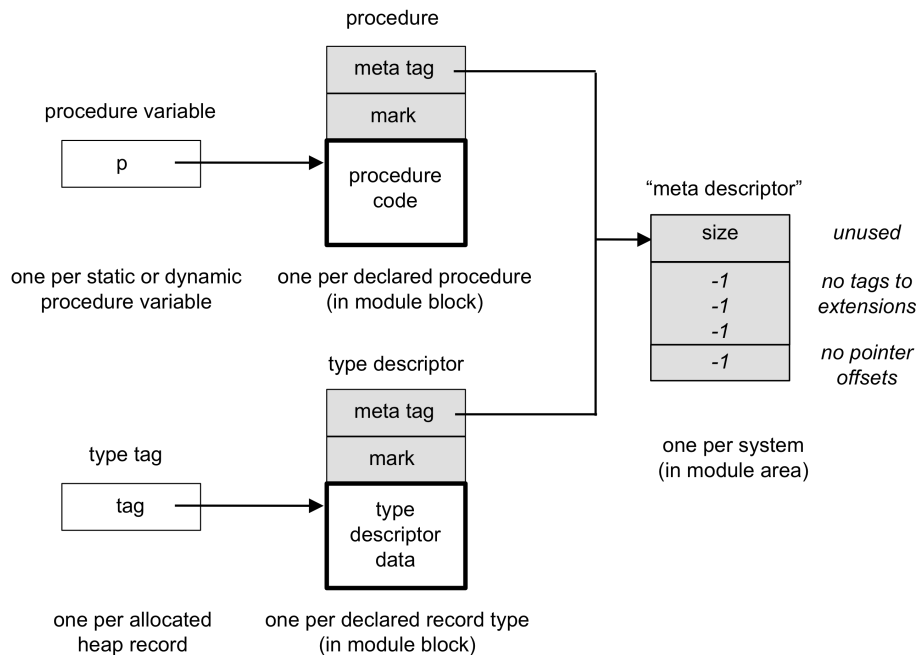


Figure 3: Procedure variable and type tag interpreted as *pointer*, and procedure and type descriptor interpreted as *record*

Alternative #3:

A simpler variant would be to treat procedure variables and type tags as *special cases* during the *mark* phase, eliminating the need for a *meta tag* field as well as the shared *meta descriptor*. The resulting run-time representation of *procedures* and *type descriptors* is shown in Figure 4.

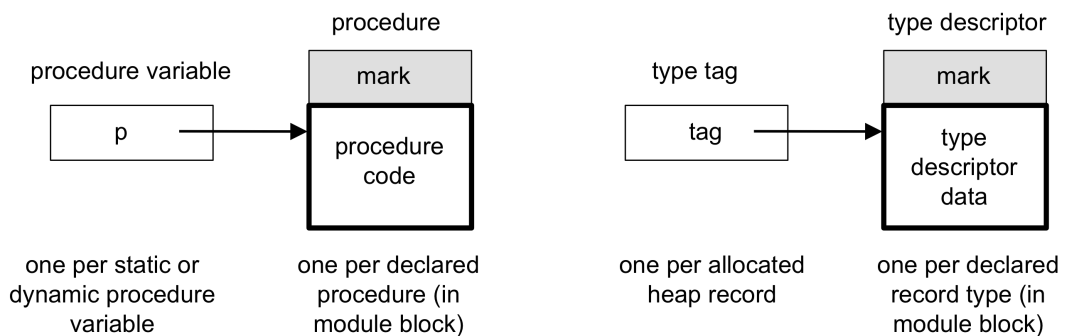


Figure 4: Procedures and type descriptors with an additional *mark* field

With these preparations, the *mark* phase of reference checking could be suitably *extended* to also include *procedure variables* in the list of “pointers” to be traversed. In sum, the mark phase would not only “touch” all reachable *heap* objects, but also those in the (static) *module blocks*.

While this technique appears appealing at first glance, a few points are worth mentioning. First, the *extended* mark phase requires an extra *mark* field to be inserted as a prefix to *each* procedure and *each* type descriptor in the module block. Given that most procedures and type descriptors are never referenced, this appears to be overkill. One could therefore decide to add

the *mark* field only to *module descriptors* rather than individual *procedures* or *type descriptors* during the *mark* phase. While this would make the check whether any of the selected modules is referenced a trivial task, it would render the *mark* phase more complex, as it would now need to *locate* the module descriptor belonging to a given procedure or type descriptor (on systems where additional meta-information is present in the run-time representation of loaded modules, such as the locations of procedures and type descriptors in module blocks or *back pointers* from each object of a module to its module descriptor, the mark phase would be simpler; however, neither FPGA Oberon 2013 nor Experimental Oberon offer such *metaprogramming* facilities).

Second, one would still need to mark *all* objects, for the same reason as outlined above, i.e. one cannot simply exit in the middle of the *mark* phase without additional action.

Third, one would also still need an extra *scan* phase to unmark the marked portion of the heap.

Finally, a comparison of the code required to implement the various alternatives showed that our solution is *by far* the simplest: the combined implementation cost of *all* modifications to the runtime representation of type descriptors and descriptors of global module data, the object and symbol file formats and the module loader, is only about 15 source lines of code²², while the *reference checking* phase itself amounts to less than 75 lines (the *total* implementation cost to realize *safe module unloading*, including the ability to unload module *groups* and the automatic collection of no longer referenced hidden modules, amounts to about 250 source lines of code).

* * *

²² See procedures *ORB.InType*, *ORB.FindHiddenFields*, *ORG.BuildTD*, *ORG.Close* and *Modules.Load*.

Appendix: Historical notes on module unloading in the Oberon operating system

In most Oberon implementations, only *client* references are checked prior to module unloading. *Type* and *procedure variable* references are usually not checked, although various approaches are typically employed to address the case where such references do exist. These approaches can be broadly grouped into two main categories: (a) schemes that explicitly allow invalidating references, and (b) schemes where all references must remain unaffected.

a. Schemes that explicitly allow invalidating references

In such schemes, unloading a module from memory *may*, and in general *will*, lead to “dangling” *type* and *procedure variable* references pointing to module data that is no longer valid²³. This includes the important case where a structure rooted in a variable of base type T declared in a base module M contains elements of an extension T' defined in a client module M'. Such elements typically contain both *type* references (type tags) and *procedure variable* references (handler procedures) referring to module M'. This is common in the Oberon viewer system, for example, where M is module *Viewers*. If the client module M' is unloaded and the module block previously occupied by it is overwritten by a module loaded later, any still existing references to M' become *invalid* at that moment. Global procedure variables declared in other modules may also refer to procedures declared in module M', although this case is much less common (global procedure variables tend to be used mainly for procedures declared in the same module).

A variety of approaches have been employed in different implementations of the Oberon system to cope with the introduced dangling *type* or *procedure variable* references:

1. On systems that use a *memory management unit* to perform virtual memory management, such as Ceres-1 or Ceres-2, one approach is to *unmap* the module space of an unloaded module from virtual memory, thereby *invalidating* future references to it. When this is done, any still existing *type* or a *procedure variable* references point to a now *unallocated* page, and consequently any attempt to access such a page later will result in a *trap*.

We consider this an unfortunate proposal for several reasons. First, users generally have no way of knowing *whether* it is in fact safe to unload a module, yet they are allowed to do so. Second, after having unloaded it, they still don't know whether references from other loaded modules still exist – until a *trap* occurs. But then it is usually too late. While the trap itself will actually *prevent* a system crash as intended²⁴, the user may *still* need to restart the system in order to recover an environment without any “frozen” parts, for example displayed viewers that may have been opened by the unloaded module. Note also that this solution requires special hardware support, which may not be available on all systems.

2. On systems that do *not* use virtual memory, such as Ceres-3 or FPGA Oberon on RISC, the easiest way to cope with dangling references is to simply *ignore* them, i.e. to *always* release the memory associated with a module without any further precautions (unless clients exist). While an attempt to access an unloaded module goes undetected *initially*, the system is still left in an *unsafe* state. It will become *unstable* the very moment another module loaded later *overwrites* the previously released module block *and* other (loaded) modules still refer to its *type descriptors* or *procedures*. This is of course undesirable.

²³ If the programming language Oberon-2 is used, there can also be references to method tables (which however are typically allocated within type descriptors).

²⁴ For example, if the unloaded module implements a subframe type, a trap is generated if the enclosing menu viewer attempts to send a “close” message to the subframe by calling its handler.

3. Another approach, which however can be used only for *procedure variable* references, is to identify *all* such references and make them refer to a *dummy* procedure, thereby preventing a run-time error when such “fixed up” procedures are called later. Of course, this solution requires the system to *know* the locations of all procedure variables in all static and dynamic objects at run time²⁵. It was used in an earlier version of Experimental Oberon, but was later discarded, mainly because the resulting effect on the *overall* behavior of the system is essentially impossible to predict (or even detect) by the user. The fact that *some* procedure variables *somewhere* in the system no longer refer to *real* but to *dummy* procedures typically becomes “visible” only through the *absence* of some action – such as mouse tracking if the unloaded module contained a viewer handler, for instance.
4. On systems that use *indirection* for procedure calls via a so-called “link table”, the same effect can be achieved by setting the *link table entries* for all referenced *procedures* of the module to be unloaded to *dummy* entries, rather than locating and modifying each individual procedure *call* that may exist anywhere in the system.

Note that using a link table to implement indirection for procedure calls is only viable on systems that provide *efficient* hardware support for it. On such systems, an “address” of a procedure is not a real memory address, but an *index* to this translation table – which the caller consults for every procedure *call*, in order to obtain the actual memory location of the called procedure. Indirection for procedure calls via a link table was used in some of the earlier versions of the Oberon system on Ceres computers, which were based on the (now defunct) NS32000 processor. This processor featured a *call external procedure* instruction (CXP *k*, where *k* is the index of the link table entry of the called procedure), which sped up the process of calling external procedures significantly²⁶. Later versions of the NS processor, however, internally re-implemented the *same* instruction using microcode, which negatively impacted its performance²⁷. For this and a variety of other reasons, the *CXP k* instruction – and with it the *link table* – were no longer used in later versions of the Oberon system.

5. For *type* references, it is possible to determine at *compile* time, whether a module *may* lead to references from other modules at *run* time. The criteria is the following: if a module *M'* does *not* declare record types which are extensions *T'* of an imported type *T*, then records declared in *M'* *cannot* be inserted in a data structure rooted in a variable *v* of the imported type *T* – precisely *because* they are not extensions of *T* (in the Oberon programming language, an assignment *p := p'* is allowed only if the type of *p'* is the same as that of *p* or an extension of it). One *could* therefore introduce a rule that a module *M'* can be safely dispensed *only* if it does *not* declare record types, which are extensions *T'* of an imported type *T*. The flip side of such a rule, however, is that modules that actually *do* declare such types can *never* be unloaded (unless, of course, other ways to safely unload such modules are implemented). Also, procedure variable references are not covered by this rule.

Even though most of these approaches have actually been realized in various implementations of the Oberon system, we consider none of them truly satisfactory. In our view, these schemes appear to only tinker with the symptoms of a problem that would not exist, if only one adopted the rule to *disallow* the removal (from memory) of still referenced module data.

²⁵ This can be achieved by including the locations of procedure variables in the module block as follows: the offsets of global procedure variables in the module's data section are stored in a separate array in the module's meta data section, while the offsets of procedure variables in dynamically allocated records are added to the type descriptor associated with each such record.

²⁶ The use of the link table also increased code density considerably (as only 8 bits for the index instead of 32 bits for the full address were needed to address a procedure in every procedure call). In addition, the link table used by the CXP instruction allowed for an expedient linking process at load time (as there are far fewer conversions to be performed by the module loader – one for every referenced procedure instead of one for every procedure call) and also eliminated the need for a fixup list (list of the locations of all external procedure references to be fixed up by the module loader) in the object file. A disadvantage is, of course, the need for a (short) link table.

²⁷ The internal re-implementation of the CXP instruction using microcode in later versions of the NS processor followed the general industry trend of implementing only frequent, simple instructions directly with hardware, while interpreting more complex instructions using internal microcode. In general, with the advent of highly regular reduced instruction set computers (RISCs) in the mid 1980s and early 1990s, the trend towards offering microprocessors providing a smaller set of simple instructions, most of them executing in a single clock cycle, combined with fairly large banks of (fast) registers, continued (and does so to this date).

The main issue appears to be that the moment one *allows* modules to release their associated memory even when references to them still exist from other loaded modules, the resulting *dangling* references must be dealt with *somehow* in order to prevent an almost certain system *crash*. However, *fixing up* or *invalidating* references will *always* remove essential information from the system. As a result, the run-time behavior of the modified system becomes *essentially unpredictable*, as other loaded modules may *critically* depend on the removed functionality. For example, unloading a module that contains an installed handler procedure of a *contents frame* may render it impossible to *close* the enclosing *menu viewer* that contains it, thereby leading to a system with “frozen” parts. A similar problem may occur with references to *type descriptors*, if they are not persisted in memory *after* unloading their associated modules.

b. Schemes where all references must remain unaffected

The second possible interpretation of *unloading* a module consists of schemes where all references *must* remain unaffected at all times. In such schemes, module unloading can be viewed as an implicit mandate to preserve “critical module data”, as long as references to the unloaded module still exist. Various possible ways to implement this mandate exist:

1. One could of course simply exit the *unload* command with an error message, whenever such references are detected. The user, however, may then be “stuck” with modules that he can *never* unload because they are referenced by modules over which he has no explicit control.
2. But the mandate could also be fulfilled by persisting any still referenced module data to a “safe” location before unloading the associated module. How can this be achieved?

For *type descriptors* referenced by *type tags* in dynamic records an easy solution exists: allocate them *outside* the module blocks in order to persist them beyond the lifetime of their associated module. One possibility is to allocate them in the *heap* at module load time²⁸. This has been implemented in Ceres-Oberon, for instance. Note that this approach eliminates dangling *type* references altogether and therefore also the *need* to check for them at run time – as type descriptors are now unaffected by module unloading.

For *procedures* referenced by *procedure variables* in either static or dynamic objects no such simple solution exists. If one doesn’t want to invalidate or fix up procedure variable references (as outlined in the previous section), the only way to “persist” procedures is to persist the *entire* module (recall that procedures may also access global module data or *call* other procedures declared in the same module).

We conclude that if one wants to address both type *and* procedure variable references, one *cannot* unload the module block from memory, as long as references to it still exist²⁹.

3. A straightforward way to automatically persist both type descriptors *and* procedures consists of *never* releasing the module block of a module once it has been loaded. Instead, when the user requests the unloading of a module, it is only removed from the *list* of loaded modules. This amounts to *renaming* the module, with the implication that a newer version of the same module (with the same name) can be reloaded again³⁰. This approach has been chosen in

²⁸ Note that one cannot simply move type descriptors around in memory, as their addresses are (typically) used to implement type tests and type guards. By allocating them in the heap at module load time, one avoids the need to move them to a different location when a module is unloaded.

²⁹ Of course, “mixed” variants are also possible. For example, one could allocate type descriptors in the heap at module load time (as in Ceres-Oberon), and either fix up all procedure variable references or prevent the release of a module block if such references still exist; however, most modules referenced by type tags are *also* referenced by procedure variables – this is in fact the typical case for dynamic records containing installed handler procedures. Thus, it seems more natural to employ the *same* approach for both type and procedure variable references.

³⁰ In a specific implementation, one might choose to make the module completely anonymous or modify the module name such that one can no longer import it (e.g., by inserting an asterisk).

MacOberon³¹, for instance. Since the associated memory of a module is never released, the issue of dangling type or procedure variable references is avoided altogether, as they simply cannot exist. However, it can also lead to higher-than-necessary memory usage if a module is repeatedly loaded and unloaded (typical on *development* systems). Nevertheless, such an approach may be seen as adequate on *production* systems where module unloading is rare.

4. A *refinement* of the approach outlined above consists of *initially* removing a module from the list of loaded modules (as in MacOberon), but *in addition* releasing its associated memory *as soon as* there are no more type or procedure variable references to it. If this is done in an automatic fashion (for example as part of a background process), module data is truly “kept in memory for exactly as long as necessary and removed from it as soon as possible”.

This is the approach chosen in Experimental Oberon. A variation of it was used in some versions of SparcOberon³².

In sum, module unloading schemes where references remain *unaffected* at all times avoid many of the complications that are inherent in schemes that explicitly *allow* invalidating references. A (small) price to pay is to keep modules loaded in memory, as long as references to them exist.

However, on *production* systems, there is typically no need to keep multiple copies of the same module in memory, while on *development* systems it is totally acceptable.

Finally, we note that on modern computers the amount of available memory, and therefore also the amount of *dynamic* data that may be allocated by modules, typically far exceeds the size of the module blocks holding the program code and global variables. Hence, not releasing module blocks immediately after a module *unload* operation typically has a rather negligible impact on overall memory usage.

³¹ <http://e-collection.library.ethz.ch/eserv/eth:3269/eth-3269-01.pdf> (The Implementation of MacOberon, 1990)

³² <http://e-collection.library.ethz.ch/eserv/eth:7103/eth-7103-01.pdf> (SPARC-Oberon User's Guide and Implementation, 1990/1991)