

The Extended Oberon System

Andreas Pirklbauer

5.5.2020

*Extended Oberon*¹ is a revision of the *Project Oberon 2013*² operating system and its compiler. In this document, the term *Project Oberon 2013* refers to a re-implementation of the original *Oberon* system on an FPGA development board in 2013, published at www.projectoberon.com.

Extended Oberon contains a number of enhancements, including the *Revised Oberon-2* programming language implementing a superset of the Oberon-07 language, completely *safe* unloading of modules or groups of modules, system building and maintenance tools, smooth fractional line scrolling of displayed texts with variable line spaces, multiple logical display areas and various other improvements. Some of these enhancements are of a purely experimental nature, while others serve the explicit purpose of exploring potential future extensions.

1. Revised Oberon-2 programming language

Revised Oberon-2 is a revision of the programming language Oberon-2³. The main difference to the original is that it implements a strict superset of *Revised Oberon (Oberon-07)* as defined in 2007/2016⁴ rather than being based on the original language *Oberon* as defined in 1988/1990⁵.

Original Oberon (1988/1990)



Original Oberon-2 (1991/1993)



Revised Oberon (Oberon-07) (2007/2016)



Revised Oberon-2 (2019/2020)

Its principal new features include type-bound procedures, a dynamic heap allocation procedure for fixed-length and open arrays, a numeric case statement, exporting and importing of string constants, and no access to intermediate objects within nested scopes.

The *Revised Oberon-2* language is described in more detail in a separate document⁶.

2. Safe module unloading

Most implementations of the Oberon system only check for *client* references prior to module unloading. Other types of references are not checked, although various approaches are usually employed to alleviate the situation where such references do exist. Not checking all references leaves a system in an *unsafe* state, which may become *unstable* the moment another module loaded later overwrites a previously released module block. In addition, the unloading of *groups* of modules with (potentially cyclic) references only among themselves is usually not supported.

¹ <http://www.github.com/andreaspirklbauer/Oberon-extended>

² <http://www.inf.ethz.ch/personal/wirth/ProjectOberon/index.html> (Project Oberon, 2013 Edition); see also <http://www.projectoberon.com>

³ Mössenböck H., Wirth N.: The Programming Language Oberon-2. Structured Programming, 12(4):179-195, 1991

⁴ <http://www.inf.ethz.ch/personal/wirth/Oberon/Oberon07.Report.pdf> (Revision 3.5.2016)

⁵ <http://inf.ethz.ch/personal/wirth/Oberon/Oberon.Report.pdf> (Revision 1.10.1990)

⁶ <http://github.com/andreaspirklbauer/Oberon-extended/blob/master/Documentation/The-Revised-Oberon2-Programming-Language.pdf>

Extended Oberon implements completely *safe* unloading of modules and module groups by systematically checking *all* possible types of module references from anywhere in the system:

1. *Client references* exist when other loaded modules import module M. Client modules may refer by name to exported constants, types, variables or procedures declared in module M.
2. *Type references* exist when *type tags* (addresses of type descriptors) in dynamic objects reachable by other loaded modules refer to descriptors of types declared in module M.
3. *Procedure variable references* exist when procedure variables in static or dynamic objects reachable by other loaded modules refer to procedures declared in module M.
4. *Pointer variable references to static module data* exist when pointer variables in static or dynamic objects reachable by other loaded modules refer to *static* objects declared in module M. Such references are only possible by resorting to low-level facilities and should be avoided (pointer variables should point exclusively to anonymous variables allocated in the *heap* when needed during program execution). But since they *can* in theory exist and the required metadata (the locations of all pointer variables) already exists for other reasons, a specific implementation may opt to also check for them prior to module unloading.

The mechanism of *safe* module unloading, as implemented in the Extended Oberon system, is described in more detail in a separate document⁷.

3. System building and maintenance tools

A minimal version of the Oberon system *building tools*, as outlined in chapter 14 of the book *Project Oberon 2013 Edition*, has been added. They provide the necessary tools to establish the prerequisites for the *regular* Oberon startup process either on an existing *local* system or on a bare metal *target* system connected to a *host* via a data link (e.g., an RS-232 serial line).

The boot linker (procedure *ORL.Link*) links a set of object files together and generates a valid boot file from them. It can be used to either generate a *regular* boot file to be loaded onto the *boot area* of a disk or a *build-up* boot file sent over a data link to a target system. One can also include an *entire* Oberon system in a single *boot file*. Installing it on a target is similar to booting a commercial operating system in a Plug & Play fashion over the network or from a USB stick. The command *ORL.Load* loads a valid *boot file*, as generated by the command *ORL.Link*, onto the boot area of the local disk, one of the two valid *boot sources* (the other being the data link).

The tools to build an entirely new Oberon system on a bare metal *target* system connected to a *host* system via a communication link are provided by the module pair *ORC* running on the host system and *Oberon0* running on the target. Apart from actually *building* the target system, there is a variety of other Oberon-0 commands that can be remotely initiated from the host once the Oberon-0 command interpreter is running on the target system, for example commands for system inspection, loading and unloading of modules or the remote execution of commands. Finally, there are tools to modify the boot loader itself (module *BootLoad*), a small *standalone* program permanently resident in the computer's read-only store.

The system building tools for Project Oberon 2013 and Extended Oberon are described in more detail in a separate document⁸.

⁷ <http://github.com/andreaspirklbauer/Oberon-extended/blob/master/Documentation/Safe-module-unloading-in-Oberon.pdf>

⁸ <http://github.com/andreaspirklbauer/Oberon-extended/blob/master/Documentation/The-Oberon-system-building-tools.pdf>

4. Smooth scrolling of displayed texts with variable line spaces

Extended Oberon enables completely smooth, fractional line scrolling of displayed texts with variable line spaces and dragging of entire viewers with continuous content refresh. Both *far* (positional) scrolling and *near* (pixel-based) scrolling are realized. The system automatically switches back and forth between the two scrolling modes based on the horizontal position of the mouse pointer. For the purist, such a feature may represent an “unnecessary embellishment” of Oberon, but it is simply indispensable if the system is to support touch display devices where a mouse is absent and viewers may not have scrollbars. In such an environment, continuous scrolling is the only acceptable way to scroll and presents a more natural user interface. As a welcome side effect, the initial learning curve for users new to the Oberon system and its user interface is *considerably* reduced.

5. Multiple logical display areas (“virtual displays”)

The Oberon system was designed to operate on a *single* abstract logical display area which is decomposed into a number of vertical *tracks*, each of which is further decomposed into a number of horizontal *viewers*. Extended Oberon adds the ability to create *several* such display areas (or *displays* for short) on the fly and to seamlessly switch between them. Thus, the extended conceptual hierarchy of the display system consists of the triplet (*display*, *track*, *viewer*) and consequently the underlying base module *Viewers* exports procedures to add and remove *displays*, open and close *tracks* within existing displays and open and close individual *viewers* within tracks. There are no restrictions on the number of displays, tracks or viewers that can be created. Focus viewers and text selections are separately defined for each display, but there is only one system-wide log.

This scheme naturally maps to systems with multiple *physical* monitors attached to the same computer. It can also be used to realize super-fast context switching, for example in response to a swipe gesture on a touch display device.

The command *System.OpenDisplay name* opens a new logical display with the specified name, *System.CloseDisplay id* closes an existing one. *System.ShowDisplays* lists all open displays, *System.ThisDisplay* shows the display *id* and *name* of the current display, *System.SetDisplay id* switches to a new display, *System.SetDisplayName id name* assigns a new name to an existing display, and *System.PrevDisplay* and *System.NextDisplay* “rotate” through the open displays.

The additional commands *System.Expand*, *System.Spread* and *System.Clone* are displayed in the title bar of every menu viewer. *System.Expand* expands the viewer *as much as possible* by reducing *all* other viewers in its track to their minimum heights, leaving just their title bars visible. The user can switch back to any of the minimized viewers by clicking on *System.Expand* again in any of these title bars. *System.Spread* evenly redistributes all viewers vertically. This may be useful after having invoked *System.Expand*. *System.Clone* opens a new logical display on the fly *and* displays a copy of the initiating viewer *there*. The user can then toggle between the two copies of the viewer (i.e. switch logical *displays*) with a single mouse click.

6. Various other improvements

Various improvements to the Oberon compiler and the runtime system have been introduced. For example, in the *Project Oberon 2013* system, the following code is generated for computing the address of an *external* global variable:

LDR RH, MT, mno*4
ADD RH, RH, offset

base address of data section of the imported module
offset computed by the loader from object's export number

where the base address is fetched from a table global to the entire system. This table contains one entry for every module loaded, namely the address of the module's data section (called the module's *static base*). The address of this table is permanently held in register MT (=R12). We note that this instruction sequence imposes a restriction of 64KB for the offset from the static base of an *imported* module, as the ADD instruction has an operand field of 16 bits only.

The same instruction sequence is generated for computing the address of a global variable of the *same* module, if the offset from the module's static base is *less* than 64KB. If, however, this offset is greater than 64KB, the following four-instruction sequence is generated instead:

LDR RH, MT, mno*4
MOV1+U RH+1, offset DIV 10000H
IOR RH+1, RH+1, offset MOD 10000H
ADD RH, RH, RH+1

base address of the module's data section (static base)
upper 16 bits of the offset from the static base
lower 16 bits of the offset from the static base
RH := static base + offset from the static base

In sum, Project Oberon 2013 uses a global module table, whose address is permanently held in a special register MT, and the base address of a module's data section is held in a register at run time. Offsets used as instruction operands are always *relative* to this base address.

By contrast, Extended Oberon uses neither a global module table nor a register holding the base address of the data section of a module. Instead, all global and external data references are coded in a special format in the module's object file and are replaced by *absolute* addresses at module load time⁹, always leading to the following two-instruction sequence at run time:

MOV1+U RH, (staticbase + offset) DIV 10000H
IOR RH, RH, (staticbase + offset) MOD 10000H

upper 16 bits of (staticbase + offset)
lower 16 bits of (staticbase + offset)

If the value of the accessed variable is to be loaded into a register, the above IOR instruction is replaced by an LDR instruction.

This approach has several advantages. First, accessing a global variable of the *same* module is no different than accessing an *external* variable. For accesses to data elements within the same module at offsets greater than 64KB, this leads to higher code density and faster code. Second, it eliminates the need for a module table global to the entire system, and with it the need for a dedicated MT register permanently pointing to this table. Third, there is no 64KB restriction on the offset from the base address of the data section of a module. Finally, the concept of a static base *register* holding the address of a module's data section is also eliminated. In earlier versions of Project Oberon 2013, this base address was held in a dedicated register SB (=R13), in the current version it can be held in any register RH. In Extended Oberon, the address of a module's data section is only used during loading/linking, but not at program execution time.

* * *

⁹ If a data element of the *same* module is accessed and the offset from the base address of the module's data section is greater than 64KB (requiring more than 16 bits), the offset is initially spread among the two-instruction sequence during compilation (8 bits in the first, 16 bits in the second) and reassembled (and subsequently converted to an absolute address) by the loader.